

Lance:

Priority-Driven Data Storage and Extraction for Sensor Networks

Geoff Werner-Allen and Matt Welsh

Harvard University

Stephen Dawson-Haggerty

UC Berkeley

Omar Marcillo and Jeff Johnson

UNH



The Problem

Sensor networks increasingly used for *data intensive* applications:

- Seismic and acoustic monitoring of volcanoes, fault zones, bridges, and buildings
- Vibration monitoring of industrial equipment
- Acoustic monitoring of animal habitats

Data rates exceed limited bandwidth and storage resources:

- Typical sensor nodes exhibit ~ 100 Kbps radio bandwidth and ~ 1 MB flash
- Best reliable transfer protocol exhibits < 8 Kbps over multiple radio hops
- Yet, each sensor node may be sampling multiple channels at 100s to 1000s of Hz

Applications must direct limited resources to extracting the *“most interesting”* signals from the network

- Of course, “most interesting” depends on the app needs and the resources available

The Solution: Lance

- System for **priority driven** allocation of bandwidth and storage resources within a sensor network
 - Sensor nodes assign **priority** to data based on application specific metric
 - Priority drives allocation of nodes' flash storage
 - Data **summaries** sent to base station, used for allocating radio bandwidth for data extraction
- Lance decouples **mechanisms** for resource allocation from app-specific **policies**
 - App-supplied **policy modules** permit a wide range of allocation policies to be implemented
 - Can target many optimization metrics: priority maximization, fairness, spatial/temporal data distribution, and more




Talk Outline


- Motivation: Volcano monitoring
- The Lance Architecture
- Policy Modules: Application-Specific Customization in Lance
- Simulation study: Lance achieves *near-optimal* allocation efficiency
- Results from deployment at Tungurahua Volcano

Wireless Sensor “Motes”


- Tmote Sky platform (Moteiv, Inc.)
 - 8 MHz (TI MSP430) CPU, 10 KB RAM, 60 KB ROM
 - 2.4 GHz IEEE 802.15.4 (“Zigbee”) radio (Chipcon CC2420)
 - 1 MByte flash for data logging
- Designed for low power operation
 - 1.8 mA CPU active, 20 mA radio active
 - 5 uA current draw in sleep state
- Runs a lightweight embedded OS, called TinyOS (www.tinyos.net)
- Cost: About \$75 (with no sensors or packaging)




Sensor Networks for Volcanic Monitoring




Existing Volcanic Sensor Station




Our Wireless Volcano Monitoring Sensor Node




Original Design



Original Design



Original Design



Original Design


GPS receiver
for time sync

*Base station tells nodes to stop
sampling – to save event in flash*




FreeWave
radio modem


Base station
at observatory




Original Design



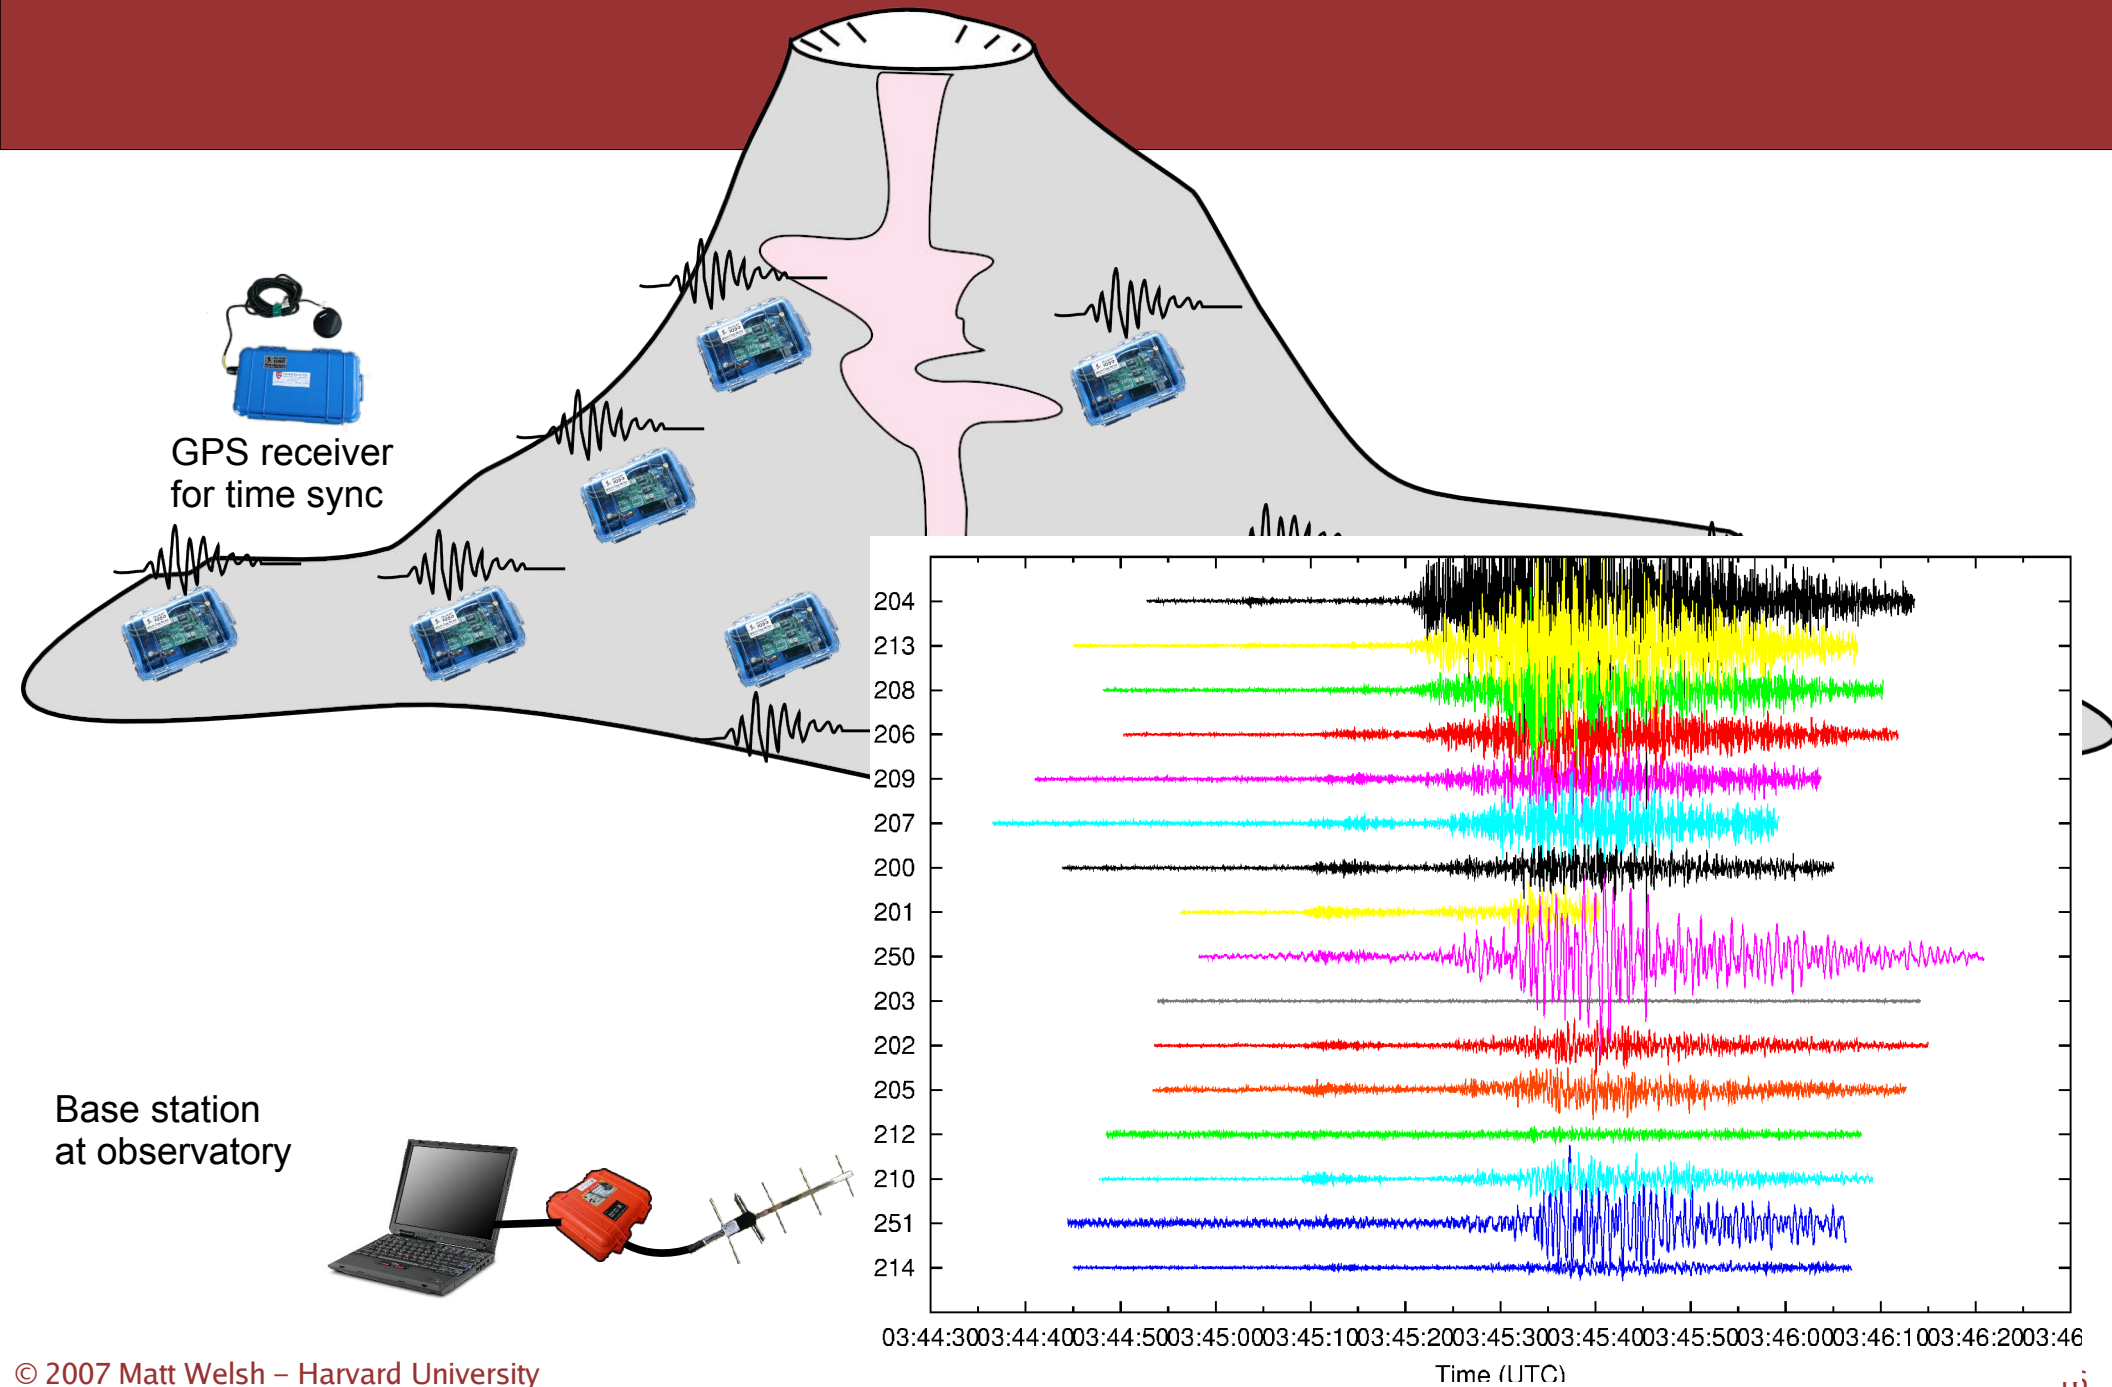
Original Design



Original Design



Original Design



Deployment at Reventador Volcano, August 2005

Next node
163m away



Radio modem

GPS receiver


Konrad

Four-channel
sensor node

Solar panels for charging
car battery (used by
FreeWave and GPS only)



Reventador Deployment Map



Reventador Deployment - Results

- 16 nodes deployed for 3 weeks in August 2006
 - Collected data on hundreds of earthquakes, eruptions, explosions.
- Lots of lessons learned [Werner-Allen et al., OSDI 2006]
 - Reliability issues with reprogramming sensors over the air
 - Time synchronization protocol bug required extensive data post-processing
 - In-depth validation of data collected compared to traditional wired data-logger
- What (else) went wrong?
 - Nodes would trigger on a small earthquake
 - Network would stop sampling and start downloading data
 - Then “the big one” would hit... and we'd fail to capture any of it!

How can we do better?

- Original system used FIFO storage and bandwidth management:
 - Treat flash as a circular buffer
 - Download one event at a time from all nodes following a trigger
 - Focused only on capturing discrete seismic events from all nodes in the network
- Instead: Use application-defined priority to drive resource allocation
 - Assign priority to each Application Data Unit (ADU) sampled by the network
 - Use priority to manage local node storage resources
 - Use priority to drive download process
- Main issues:
 - How do we compute and manage priorities?
 - How to tailor this approach for different applications?
 - How can we target different optimization metrics? (Data quality, fairness, etc.)?

Core Resource Limitations

Sensor nodes have only 1 MByte of flash

- Enough to store ~ 20 minutes of data

Reliable transfer protocols are slow!


- Best speed we've seen is around 500 bytes per second
- Often do much worse: below 100 bytes/sec in some cases
 - *Can take several minutes to download 60 sec. of data from a single node*
- Flush [Dutta et al.] claims up to 1024 bytes/sec in multihop cases
- Also, can only reasonably perform one download at a time

So, need to be careful about how we allocate storage and bandwidth.

Defining Priority

Core assumption: some data is “better” than others.


- Application's goal is to extract the “best” data given resource limitations



Defining Priority

Core assumption: some data is “better” than others.

- Application's goal is to extract the “best” data given resource limitations



Problem Definition

- Each sensor node samples ADUs $A_i = \{d_i, p_i, c_i\}$ where
 - d_i is the data, p_i is the application assigned priority, and c_i is the cost
 - Cost captures bandwidth and/or energy to reliably download Adu from the network
 - Priorities define partial order on all ADUs: $A_i \geq A_j$ iff $p_i \geq p_j$
- High-level goal: Download highest-priority ADUs such that the total cost for data retrieval is less than capacity C
 - Where capacity expressed in terms of bandwidth or energy availability
- We define the *optimal set* of ADUs Ω :
 - Set of ADUs, rank-ordered by decreasing priority $\{A_1 \geq A_2 \geq \dots \geq A_k\}$ s.t. total cost is less than capacity C

The overall system goal is to download the ADUs in the optimal set Ω .

Problem Definition

Determining Ω requires complete knowledge of all ADUs over all time.

- When determining whether to download a given ADU, must know whether it is in Ω .
- No way of knowing this without knowledge of the future ADUs that will be sampled.

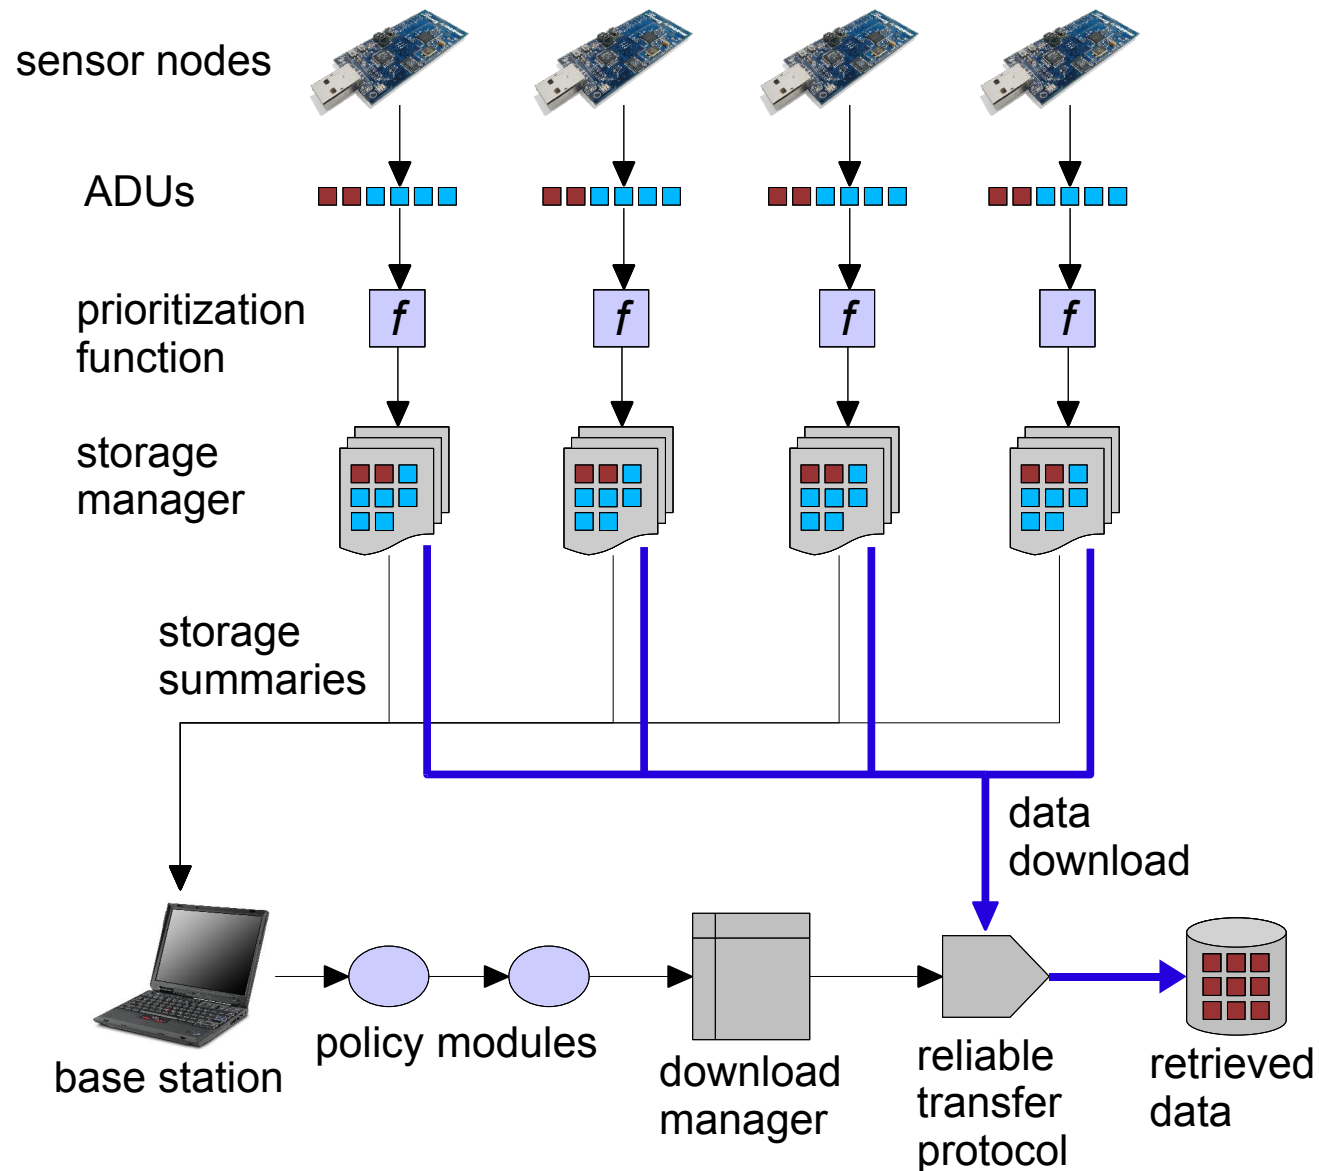
Also, retrieving Ω may require infinite per-node storage.

So, we must come up with an *online* algorithm that can operate with *limited* per-node storage capacity.

Efficiency metric: Coverage of downloaded set S with Ω

- Coverage = $|S \cap \Omega| / |\Omega|$
- Note that this metric gives no credit to ADUs downloaded *not* in the optimal set Ω
- Other metrics possible too ... more later.

Lance Architecture



Lance Storage Manager

- Nodes compute priority for each sampled ADU
 - Prioritization function is app-specific
 - Must not consume inordinate resources
- Nodes treat local flash as bounded-size priority queue
 - When storing new ADU, evict lowest-priority ADU first
- Flash technology imposes some limitations:
 - Must erase entire sector (e.g., 64 KB) before new data can be written
 - *We match size of ADU to one sector: About 109 sec. of data at 100 Hz*
 - Sector erase is slow – 0.6 sec typical for ST M25P80
 - Therefore, we perform sector erase concurrently with storing previous ADU
 - *Must evict a sector before we know what will be stored in it!*

Lance Download Manager

- Design goal: Perform bandwidth allocation *centrally*, at the base station
 - Simpler, and far more robust, than a complex decentralized algorithm.
 - Allows network's behavior to be radically changed by tweaking policies at base, without reprogramming nodes.
- Nodes send periodic storage summaries to the base station
 - List of ADUs stored on each node with corresponding priority and timestamp
- Storage summaries used by download manager to assign *download priority* for each ADU
 - Recall: We perform one ADU download at a time to avoid network congestion

Download Manager Policy

Simplest policy: Download priority == node assigned storage priority

- Simply rank-order all ADUs stored in the network by priority
- Download the highest-priority ADU

Problem: Download priority for an ADU may depend on *global* network state


- Example: Achieving fairness across nodes, or *normalizing* ADU priorities across nodes
- *Correlated event detection*: Detect spike in priority across several nodes and download same time window from all nodes

Individual nodes may not be aware of the network-wide priority for a given ADU

- Want to allow application to modify the “raw” ADU priorities to implement a wide range of download policies.

Lance Policy Modules

- User-supplied functions to inspect, filter, or modify raw ADU priorities at the base station.
- Composed into a linear chain:




- Take stream of ADU priorities as input, emit (possibly modified) priorities as output
- Can maintain internal state
- Must run efficiently (i.e., keep up with stream of ADU priorities from the network).

Example Policy Modules

- Priority thresholding
 - **filter**: Set priority to 0 if input priority below threshold T
 - *Download manager will not download an ADU with zero priority value*
- Noise removal and calibration
 - **adjust**: adjust raw priority by adding or subtracting fixed offset
 - **debias**: normalize priority values across nodes
- Priority dilation
 - **timespread**: assign high ADU priority values to ADUs adjacent in time
 - **spacespread**: assign high ADU priority values to ADUs sampled by different nodes
- Cost weighting
 - **costweight**: scale ADU priority by cost to download
 - *e.g., based on number of radio hops from the base station*

Example: Correlated Event Detection

- **correlated** policy module $W(k, \delta)$
 - counts number of ADUs within a time window δ with a nonzero priority value
 - if at least k ADUs match, retain input priorities for ADUs in the window
 - otherwise, set ADU priorities in window to 0
- Implementing our original monitoring system in Lance:



- Policy now implemented at the base station, rather than on motes.
- Easy to modify behavior of network just by changing policy modules.

Evaluation Metrics

Recall: Definition of optimal set Ω

- Set of ADUs that a perfect system with full knowledge of future data arrival would have downloaded, given the same constraints on bandwidth and storage.

Coverage metric: $|S \cap \Omega| / |\Omega|$

- Given downloaded set S and optimal set Ω
- Problem: doesn't matter *which* ADUs in Ω we manage to download.

We define *weighted coverage* $K(S, \Omega)$ as follows:

- Assign a score σ to each ADU as: $|\Omega| - \text{rank of ADU in } \Omega$
(Top ranked ADU gets score N , second-ranked $N-1$, etc.)

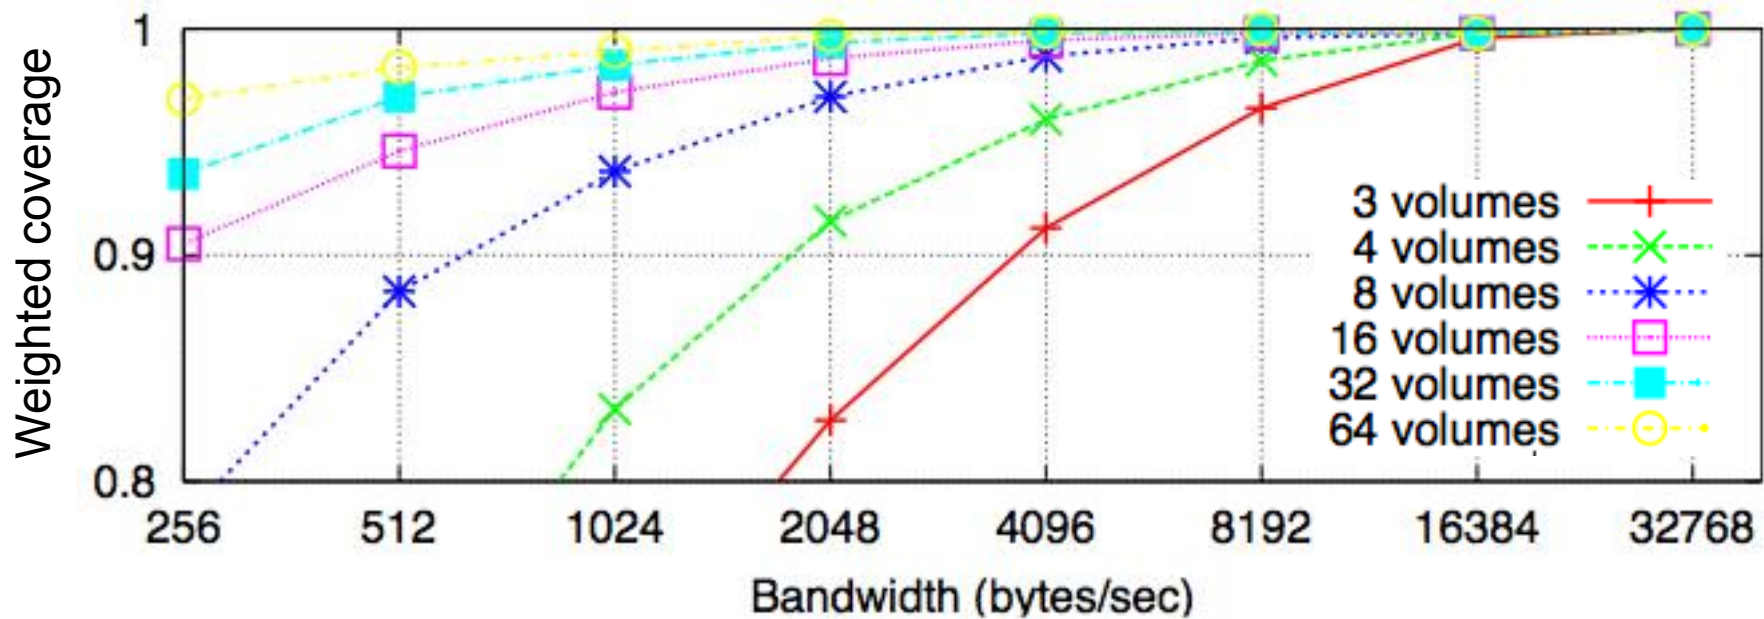
$$K(S, \Omega) = \frac{\sum_{x \in S \cap \Omega} \sigma(x)}{\sum_{y \in \Omega} \sigma(y)}$$

Methodology

- Simulated 16 sensor nodes in a radial or linear topology
 - ADU size of 64 KB and data generation rate of one ADU per node per minute
- ADU priorities drawn from three distributions:
 - Uniformly random
 - Zipf ($\alpha = 1$)
 - “Bursty Zipf”: With probability P , select a new value from a Zipf distribution; with probability $1-P$ use previous value.
 - *P controls “burstiness”: $P=1$ is equivalent to Zipf. $P=0$ all values are the same.*
- Also make use of real data sets from Reventador and Mt. St. Helens seismic sensors.


FIFO storage management

- Poor coverage under constrained bandwidth and storage capacity

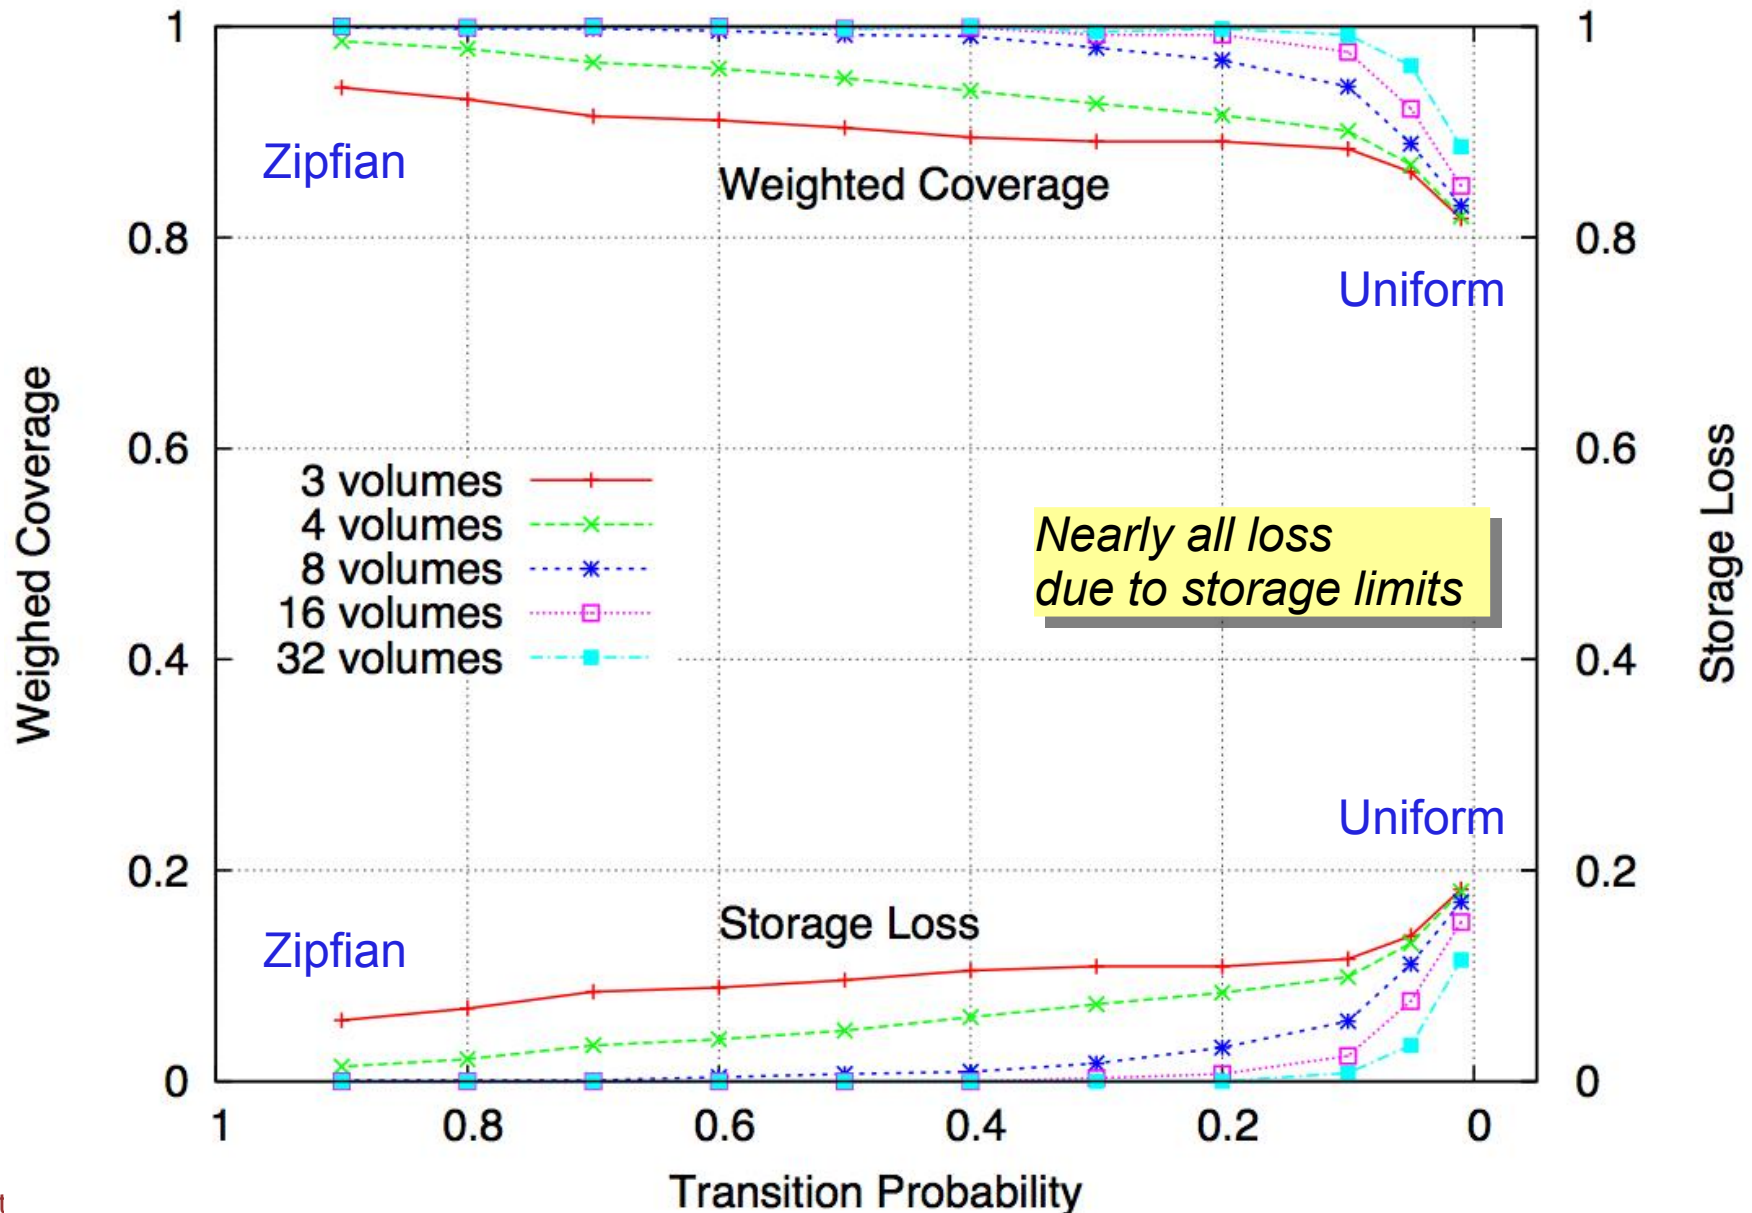


Performance using Lance

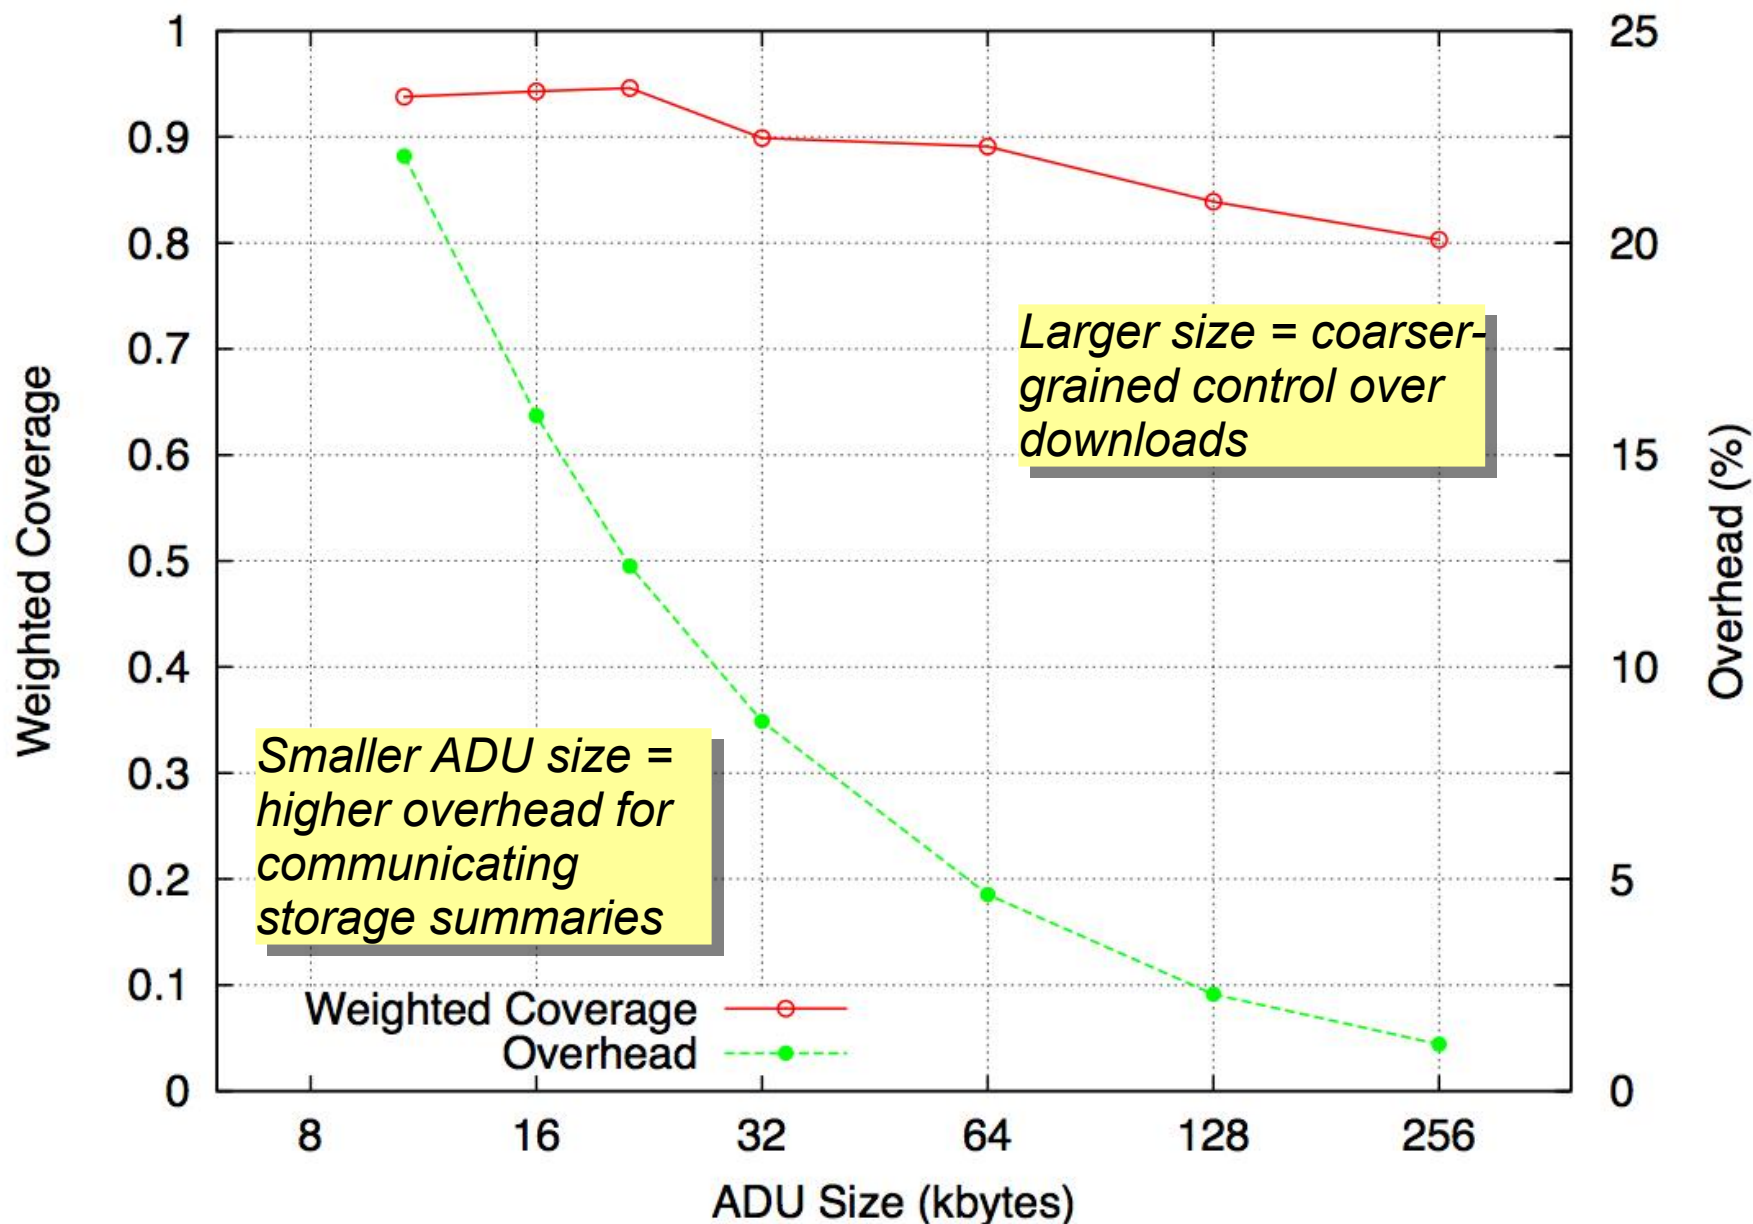
- *Near-optimal* coverage under range of resource constraints
 - Exceeds 95% in all cases




Effect of priority burstiness



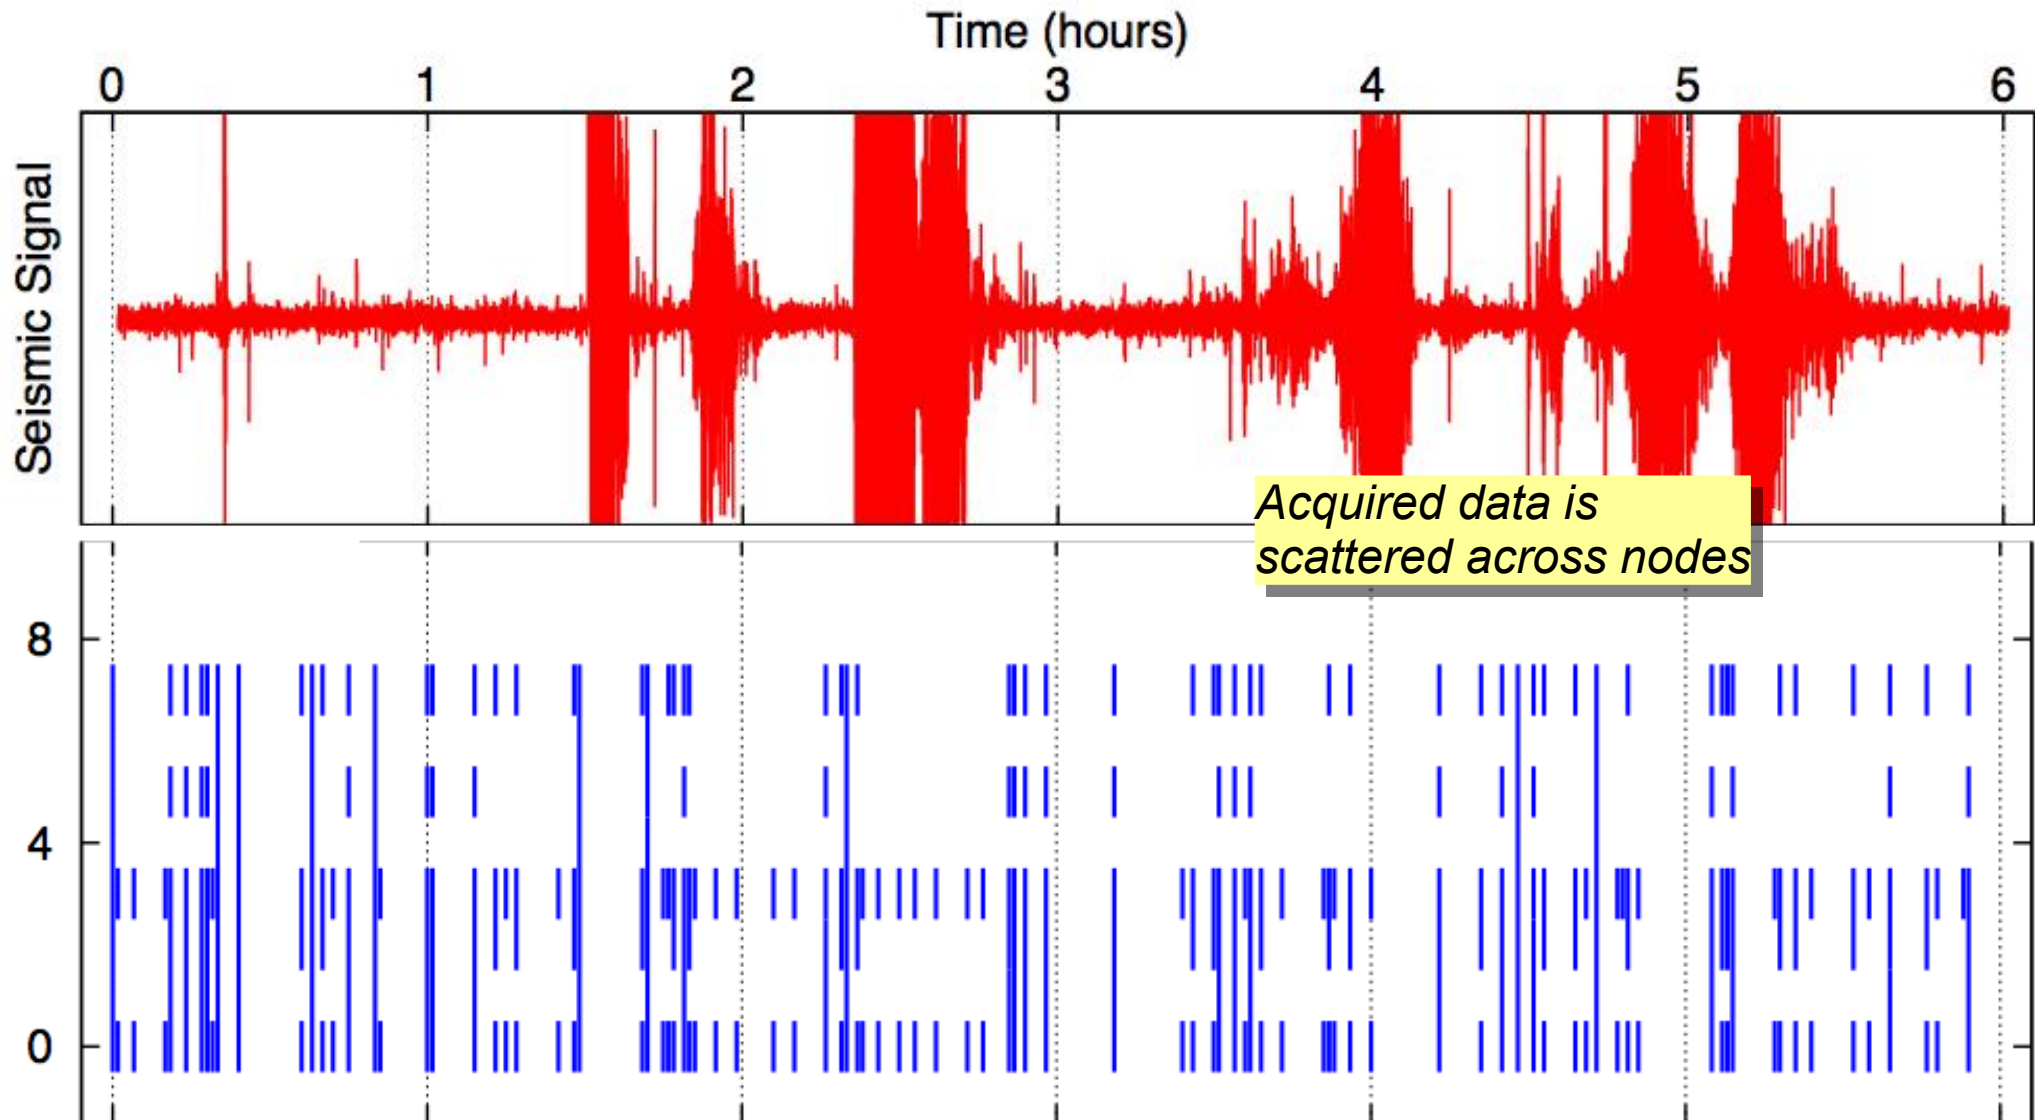
Effect of varying ADU size




Correlated event detection



Default download policy



Correlated event detection policy




Tungurahua field deployment, August 2007

Tungurahua summit

NI06
NI05
NI00
N400
NI03
NI01
NI02
NI04
750 m
Freewave

8 km to observatory








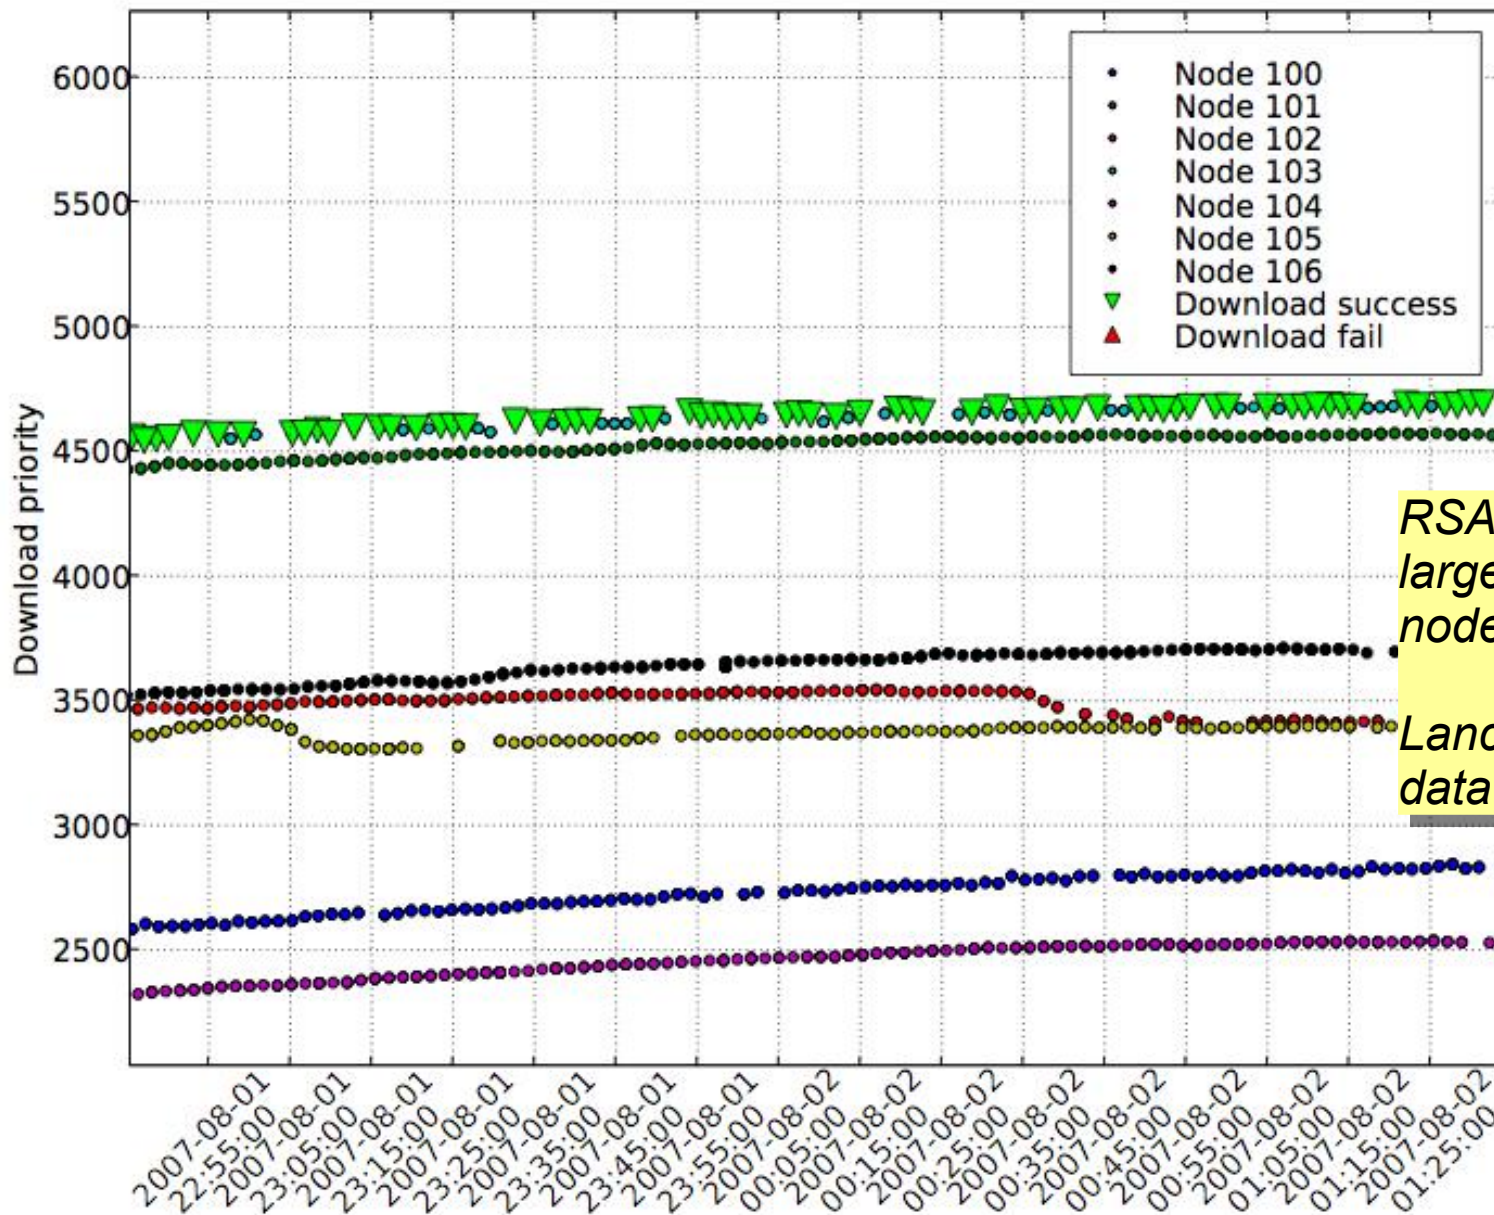
Deployment Statistics

- Ran 8 sensor nodes for a total of 71 hours
 - Lance used to manage storage and bandwidth
 - Experimented with different prioritization functions and policy modules
- Successfully downloaded 1232 ADUs (77 MB of data)
 - 308 downloads failed due to timeouts: success rate 80%
- Total storage summaries span 11012 ADUs (688 MB of data)
 - Lance downloaded 11% of the data acquired by the network
- No significant node outages observed

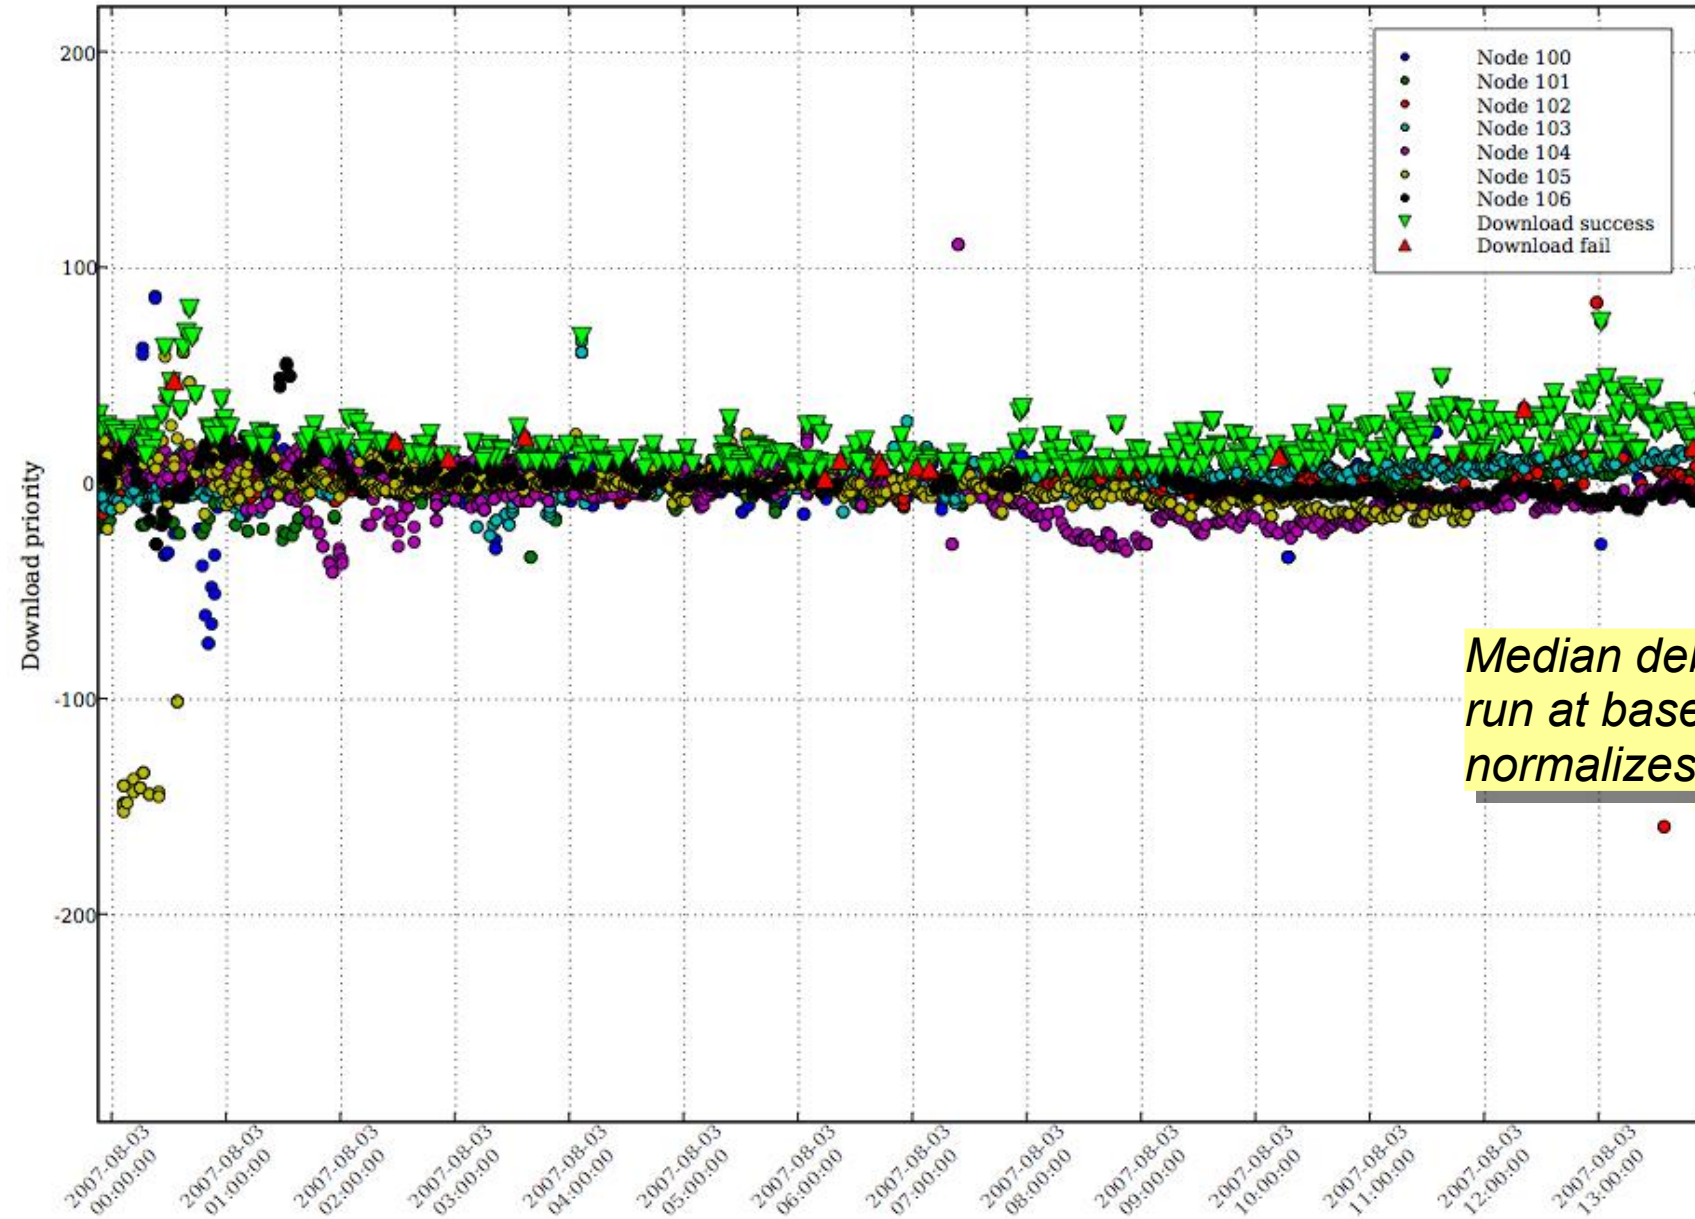
Bandwidth Breakdown



RSAM prioritization: DC bias

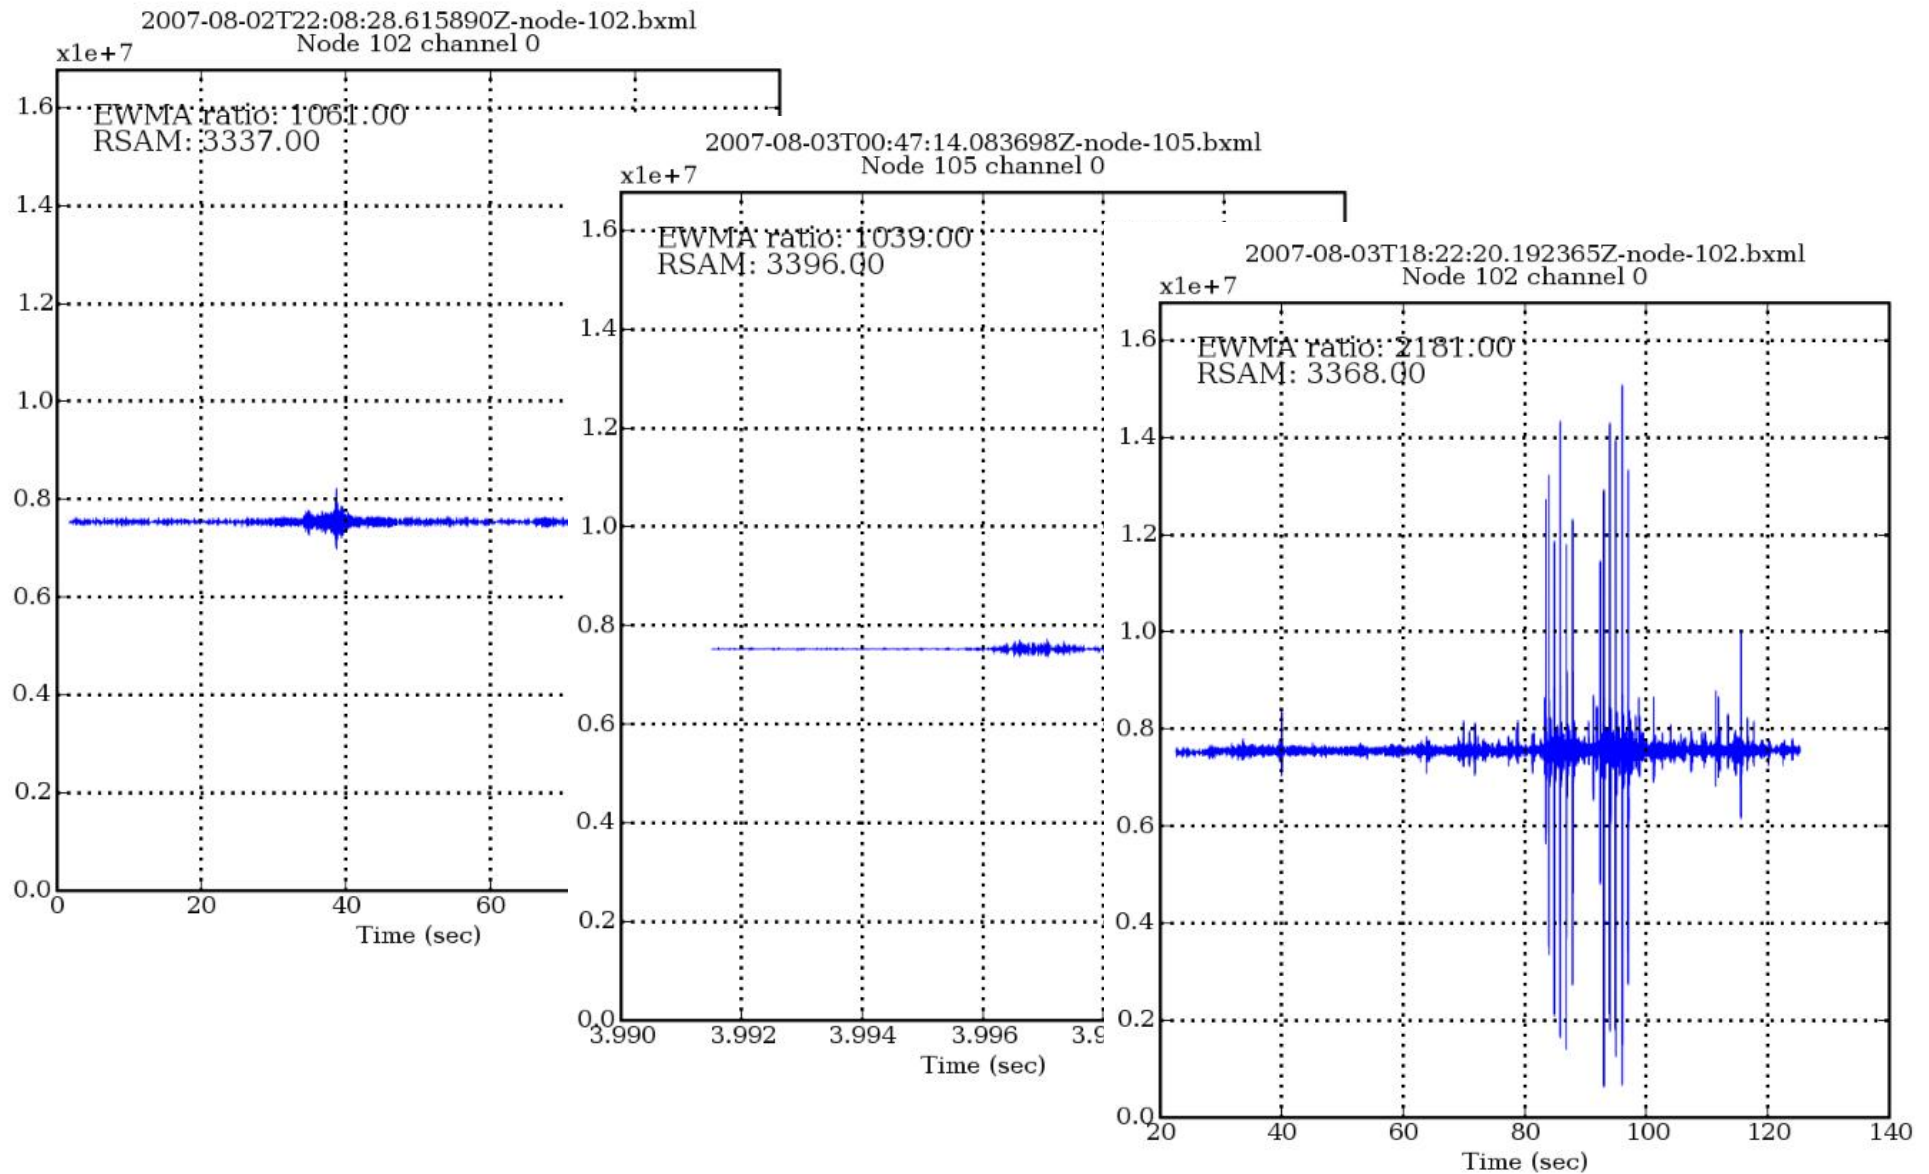


The fix: Debiasing policy module



*Median debiasing filter
run at base station
normalizes RSAM values*

We fixed the volcano...




How did RSAM perform?

- Volcanic activity was unusually low during the deployment
 - Only about 20 small earthquakes, no explosions
 - Week before, activity was much higher, with numerous explosions and dozens of earthquakes a day
- Took 8483 ADU summaries received after applying RSAM filter
 - Covering about 16 hours of the deployment
 - Computed optimal set Ω and weighted coverage $K(S, \Omega)$ for the downloaded data
- Results: Optimal set Ω included 393 ADUs. Lance downloaded 418.
 - Why more? Lance may have downloaded from nodes with faster transfer bitrates than optimal set may have chosen
 - Weighted coverage of **73%**
 - *(FIFO would have achieved only 51%)*
 - Lower than simulations, probably due to lack of variation in ADU priorities

EWMA prioritization function

- Reprogrammed nodes after 25 hours to try to trigger on earthquakes
 - Still, we observed only 9 discernible events after the reprogram.
- 11012 ADU summaries received during this time
 - Optimal system would have downloaded 554 of them.
 - Lance downloaded 518. Weighted coverage of **80%**.
 - *Fifo system would have achieved 50%*

EWMA prioritization behavior



Future Directions

- Exploring more powerful sensor node platforms
 - Xscale based platforms (e.g., iMote2) offer significant horsepower for modest energy cost
- Extend to multitier networks
 - Microservers in the field for local data collection and processing
- Move beyond data collection to in-network computation
 - Explore cost/fidelity tradeoff between raw data and extracting higher-level features
- New application domains
 - e.g., Biomedical monitoring, structural/bridge monitoring, acoustic applications
 - Limb motion analysis of patients with Parkinson's Disease (with Spaulding Rehabilitation Hospital, Boston)

Conclusions

- Wireless sensor networks can be used for data-intensive applications
 - But, radio bandwidth and storage are precious and must be managed carefully!
- Lance provides a flexible framework for maximizing network efficiency
 - Driven by application-defined *prioritization* of data
 - Node-local prioritized storage management
 - Network-wide download management
 - Policy modules enable
- Lance achieves highly efficient management of limited resources
 - Simulations: > 95% efficiency for wide range of storage capacities, bandwidths, and data distributions
 - Real deployment: efficiency of 73-80%, possibly hampered by low level of volcanic activity

Thank you!



Source code, hardware designs, and data sets available:
<http://www.eecs.harvard.edu/~mdw/proj/volcano>

Tungurahua



Banos



Image © 2006 TerraMetrics

© 2006 Google™


Pointer 1° 24'30.92" S 78° 26'06.79" W elev 7941 ft

Streaming ||||| 100%

Eye alt 18448 ft



Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
-




Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)




Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)
 - Route each chunk to base over multihop path




Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)
 - Route each chunk to base over multihop path




Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)
 - Route each chunk to base over multihop path



Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)
 - Route each chunk to base over multihop path
- Base requests missing chunks after timeout



Fetch Reliable Transfer Protocol

- Base station generates request containing:
 - node ID, block ID, bitmap of needed chunks
- Intermediate nodes flood request to network
 - Eliminates need for forward routing path from base
- Target node reads data from flash, breaks into chunks
 - One chunk per radio message (32 bytes of payload)
 - Route each chunk to base over multihop path
- Base requests missing chunks after timeout
- Node responds with missing data
-

