


# Decentralized, Adaptive Resource Allocation for Sensor Networks

Geoff Mainland, David C. Parkes, and Matt Welsh

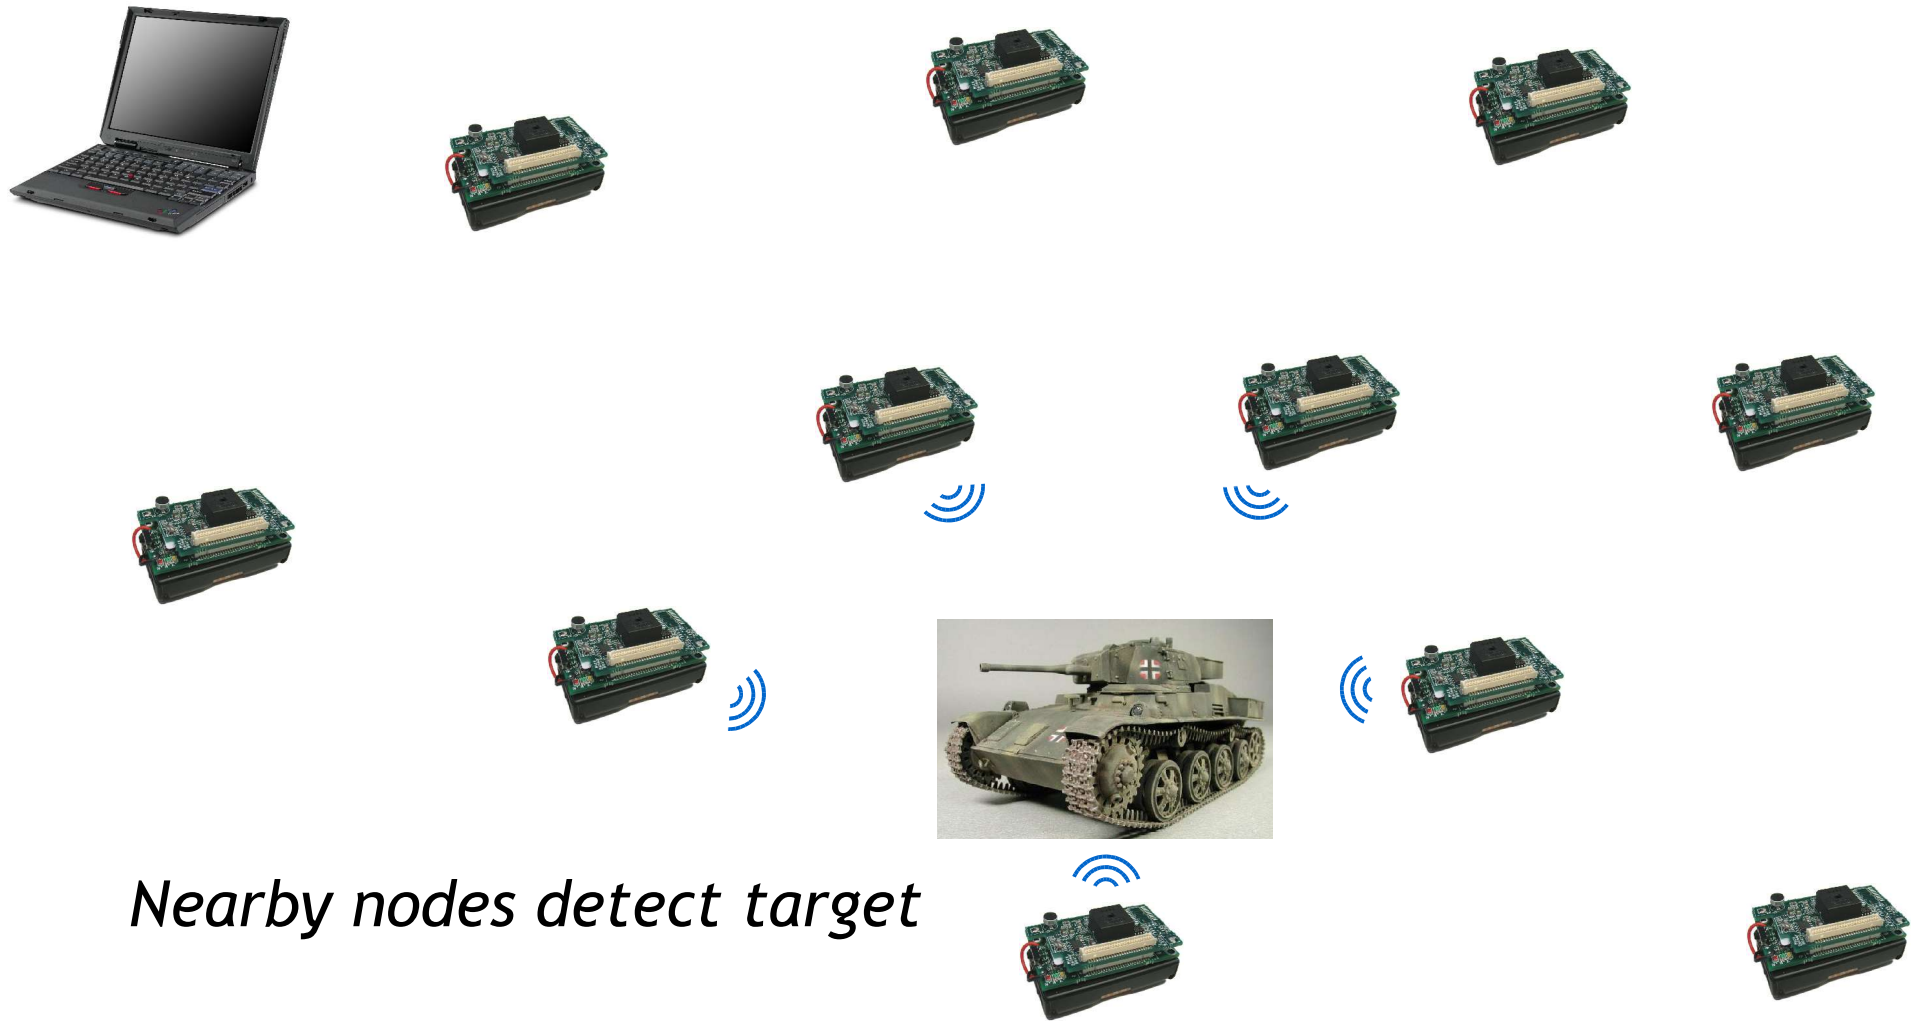
Division of Engineering and Applied Sciences

Harvard University

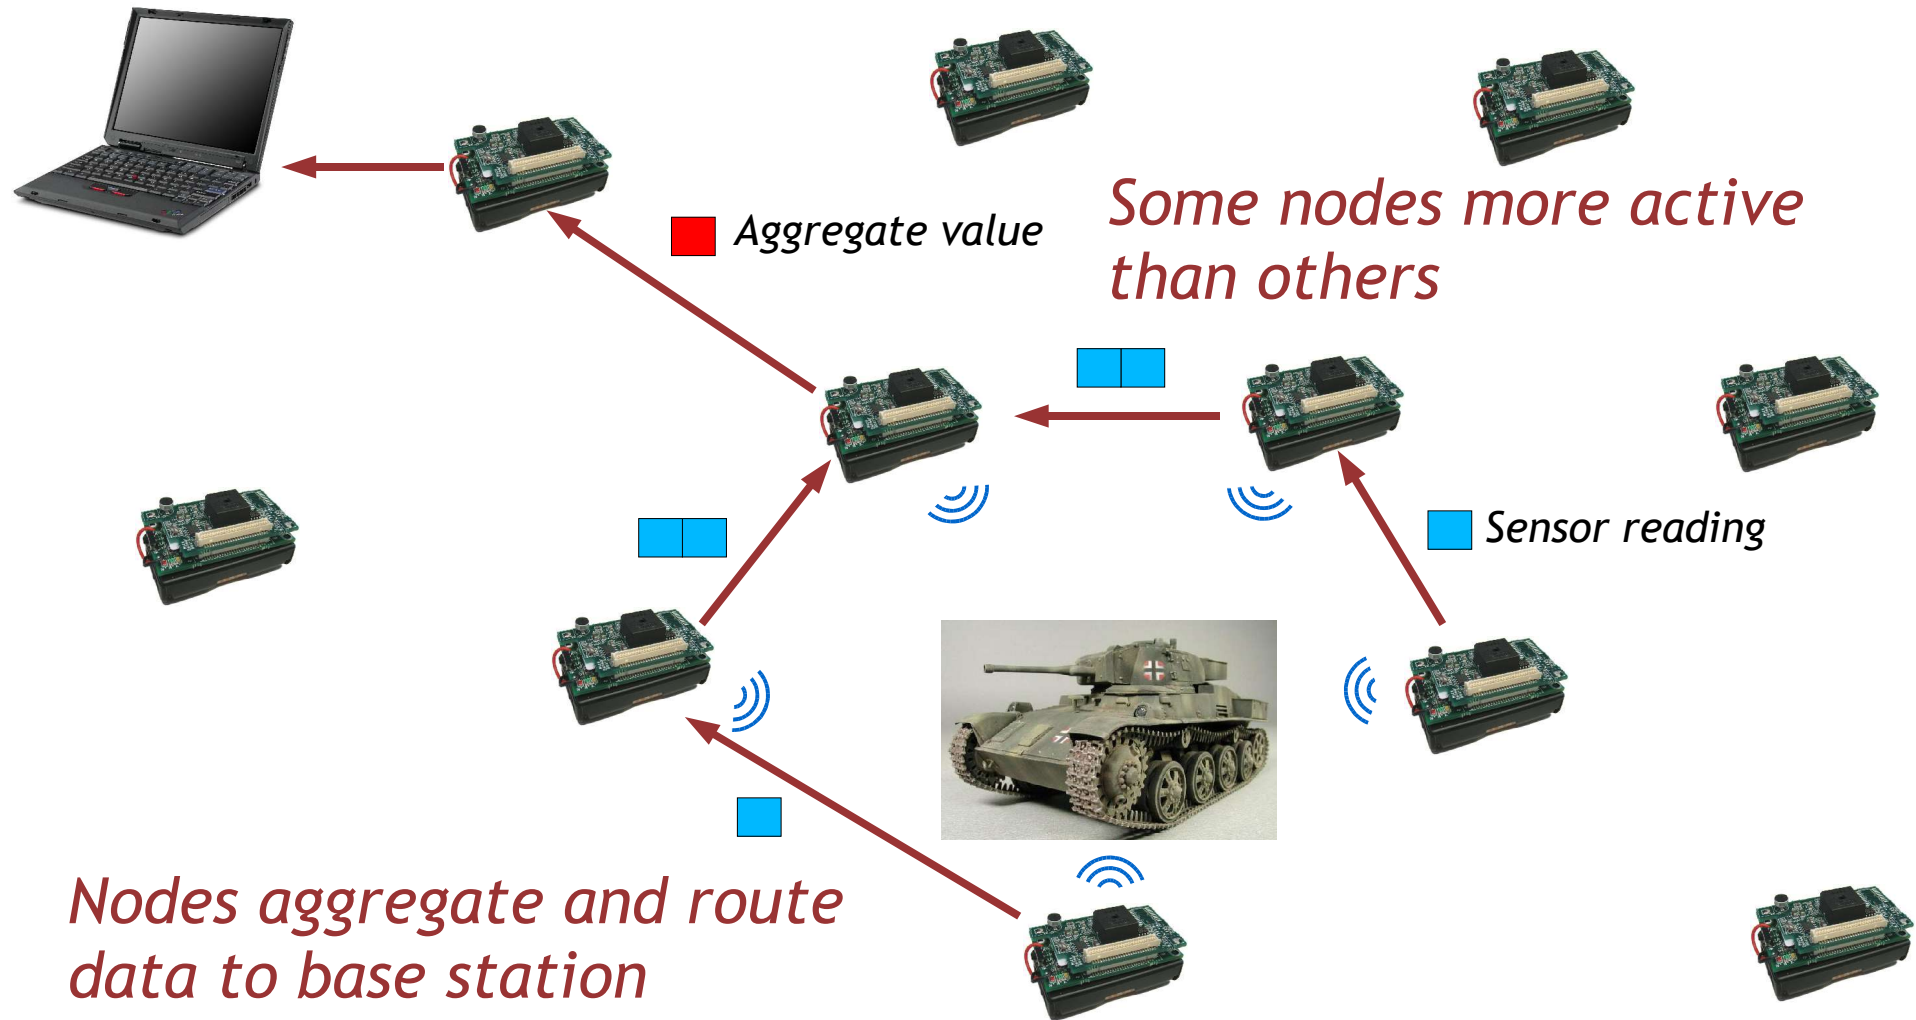
{mainland, parkes, mdw}@eecs.harvard.edu



# Motivating example: Tracking



# Motivating example: Tracking



# Problem Definition

The actions performed by each node have a deep impact on:

- **Accuracy and latency** of data returned to the base station
- **Radio contention**
- **Energy use** (and therefore, network lifetime)

Standard technique: *periodic duty cycling*

- e.g., “Every  $T$  seconds, sample the magnetometer, then transmit data if above threshold”
- Problem: Periodic sampling and communication not optimal for all nodes

Instead, nodes should *self-schedule* based on their local state

- Node should decide locally which operations to perform and how often
- Driven by interaction with environment

Node operation should *adapt* to changing conditions

- e.g., If interesting event happens nearby, node might ramp up sampling rate

# Self-Organizing Resource Allocation (SORA)

SORA is an adaptive scheduling technique for sensor networks

- Rather than static scheduling, individual nodes tune their schedules over time
- No central control or dictated node program

Goal: Nodes should avoid *wasting energy*

- Every action taken by a node consumes some amount of energy
- However, this energy is only *sometimes* “useful”

Example: Listening for incoming radio messages

- Consumes a lot of energy, but only useful if a message is actually received

Idea: Use feedback on which actions are useful to tune node behavior

- Nodes receive rewards when they take useful actions
- SORA uses reinforcement learning techniques to select best actions to take

# Self-Organizing Resource Allocation (SORA)

## Basic model:

- Nodes can select among a set of *actions*
- Each action has an associated *energy cost*
- When an action is “successful,” the node earns a *reward*

## Examples of actions (*energy cost measured on MicaZ motes*):

- *Sample* a sensor (*energy cost 84  $\mu$ J*)
- *Listen* for incoming radio messages (*energy cost 5.9 mJ*)
- *Transmit* a radio message (*energy cost 2.4 mJ*)
- *Aggregate* multiple sensor readings into a single value (*energy cost 84  $\mu$ J*)

## Each node attempts to *maximize its reward*

- That is, subject to energy constraints
- We assume that nodes can determine *locally* which actions were useful

# Utility Function

SORA drives action selection by assigning each action a *utility*

*The utility for an action is a function of:*

Reward for taking a successful action

- Advertised by base station and propagated to entire network
- Make use of lightweight data dissemination protocols (e.g., Trickle)

Energy availability

- Taking an action must stay within the node's energy budget
- We model nodes as having an *energy reserve* that is replenished at a constant rate

Data dependencies

- Cannot aggregate data until multiple samples have been received
- Cannot transmit if nothing in local buffer

# Learning Expected Rewards

The utility function is the *expected reward* for taking an action:

$$u(a) = \begin{cases} \beta(a) \times \text{reward}(a) & \text{if the action's dependencies have been met} \\ 0 & \text{otherwise (e.g., not enough energy)} \end{cases}$$

$\beta(a)$  is the estimated *probability of success* for action  $a$

- $\beta(a)$  is learned over time using an exponentially weighted moving average (EWMA)

When  $u(a)$  for all actions is zero, node performs the “sleep” action

- Places node in lowest-power state for a short period of time (0.25 sec)

Nodes explore the action space to avoid falling into local minima




- $\epsilon$ -greedy reinforcement learning
- Nodes usually take the action with highest utility  $u(a)$
- However, with (small) probability  $\epsilon$ , the node takes a *random* action



# Utility Function Example

Utility functions vary depending on node's position in network

- Nodes near target have high utility for **sampling**
- Nodes along routing path have high utilities for **listening** and **sending**



# SORA Design Issues

## Nodes operate using a very simple program

- Small amount of state and low computational overhead

## Network rapidly adapts to changing conditions

- Nodes take actions that they individually find to earn rewards
- Learning strategy for utility function adapts behavior over time

## No explicit coordination between nodes

- However, reward feedback leads to a natural equilibrium
- e.g., Fraction of nodes transmitting and listening for messages is *balanced*

## Reward prices do not have a large impact on behavior!

- They only serve to differentiate behavior when multiple actions are profitable.
- For our experiments, we set the reward for each action to the same value
- More later...


# Evaluation Methodology

## Simulation based on TinyOS environment

- Captures realistic hardware-level effects
- Target travels in circular path
- Routing using GPSR to base station



## Evaluation goals:

- Can SORA achieve good tracking accuracy?
- How efficient is the resulting resource allocation?
- How well do nodes adapt to changing conditions?




# Node Behavior over Time

(one node along the path of the vehicle)




More frequent sampling, agg, and transmission while vehicle is nearby






# Node Behavior over Time

(one node along the path of the vehicle)




*Vehicle approaches*

*Vehicle leaves*



*Occasional listening and sampling while exploring action space*



# Comparison to Alternatives

Implemented two alternative scheduling techniques

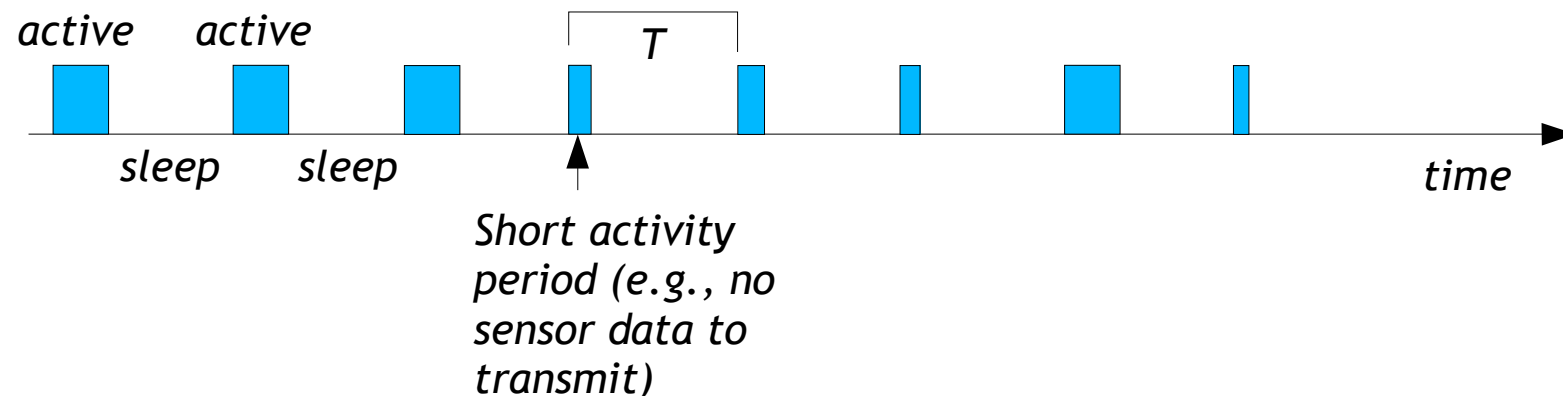
- (Plus a third described in the paper)

Each node is given a daily energy budget (e.g., 1000 J/day)

- Node's energy reserve *continually* refills at this rate

Static, periodic schedule (most commonly used technique today)

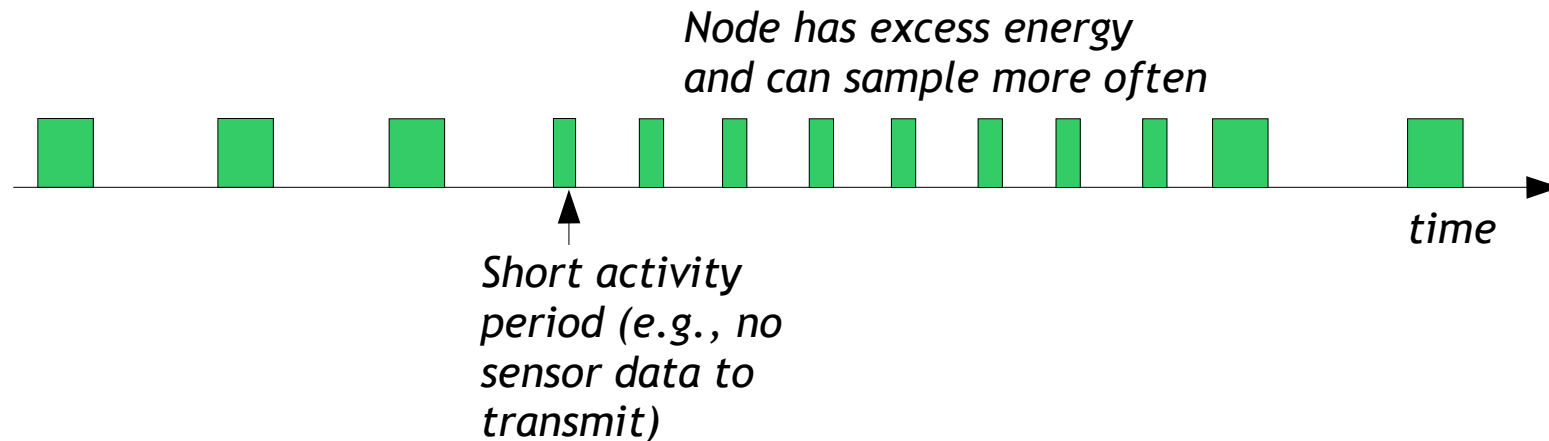
- Nodes periodically sample, listen, aggregate, transmit, and sleep
- All nodes operate at the same rate, calculated **offline** to meet energy budget
  - *This is conservative: Nodes may not use entire energy budget*



# Comparison to Alternatives

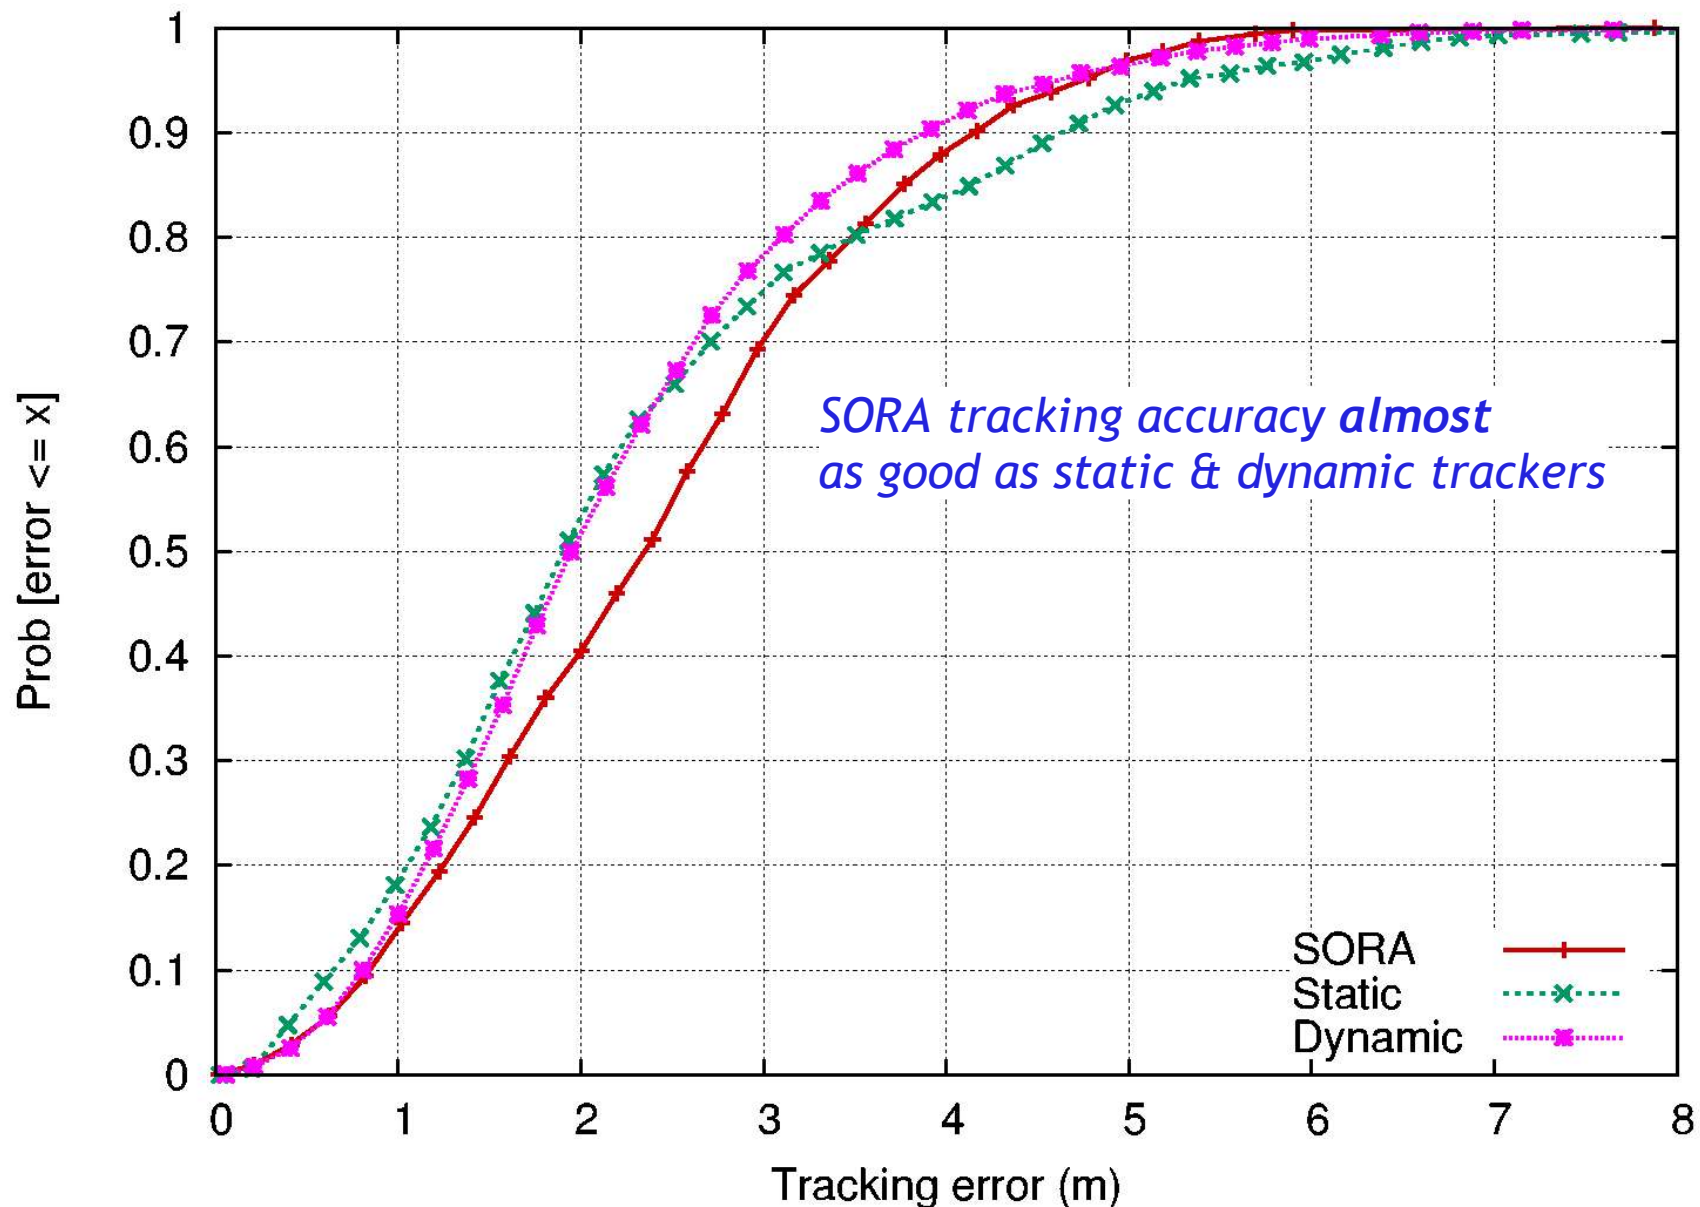
## Dynamic periodic schedule

- Nodes dynamically tune processing rate to exactly exhaust their energy reserve
- Some nodes will operate at faster rates than others



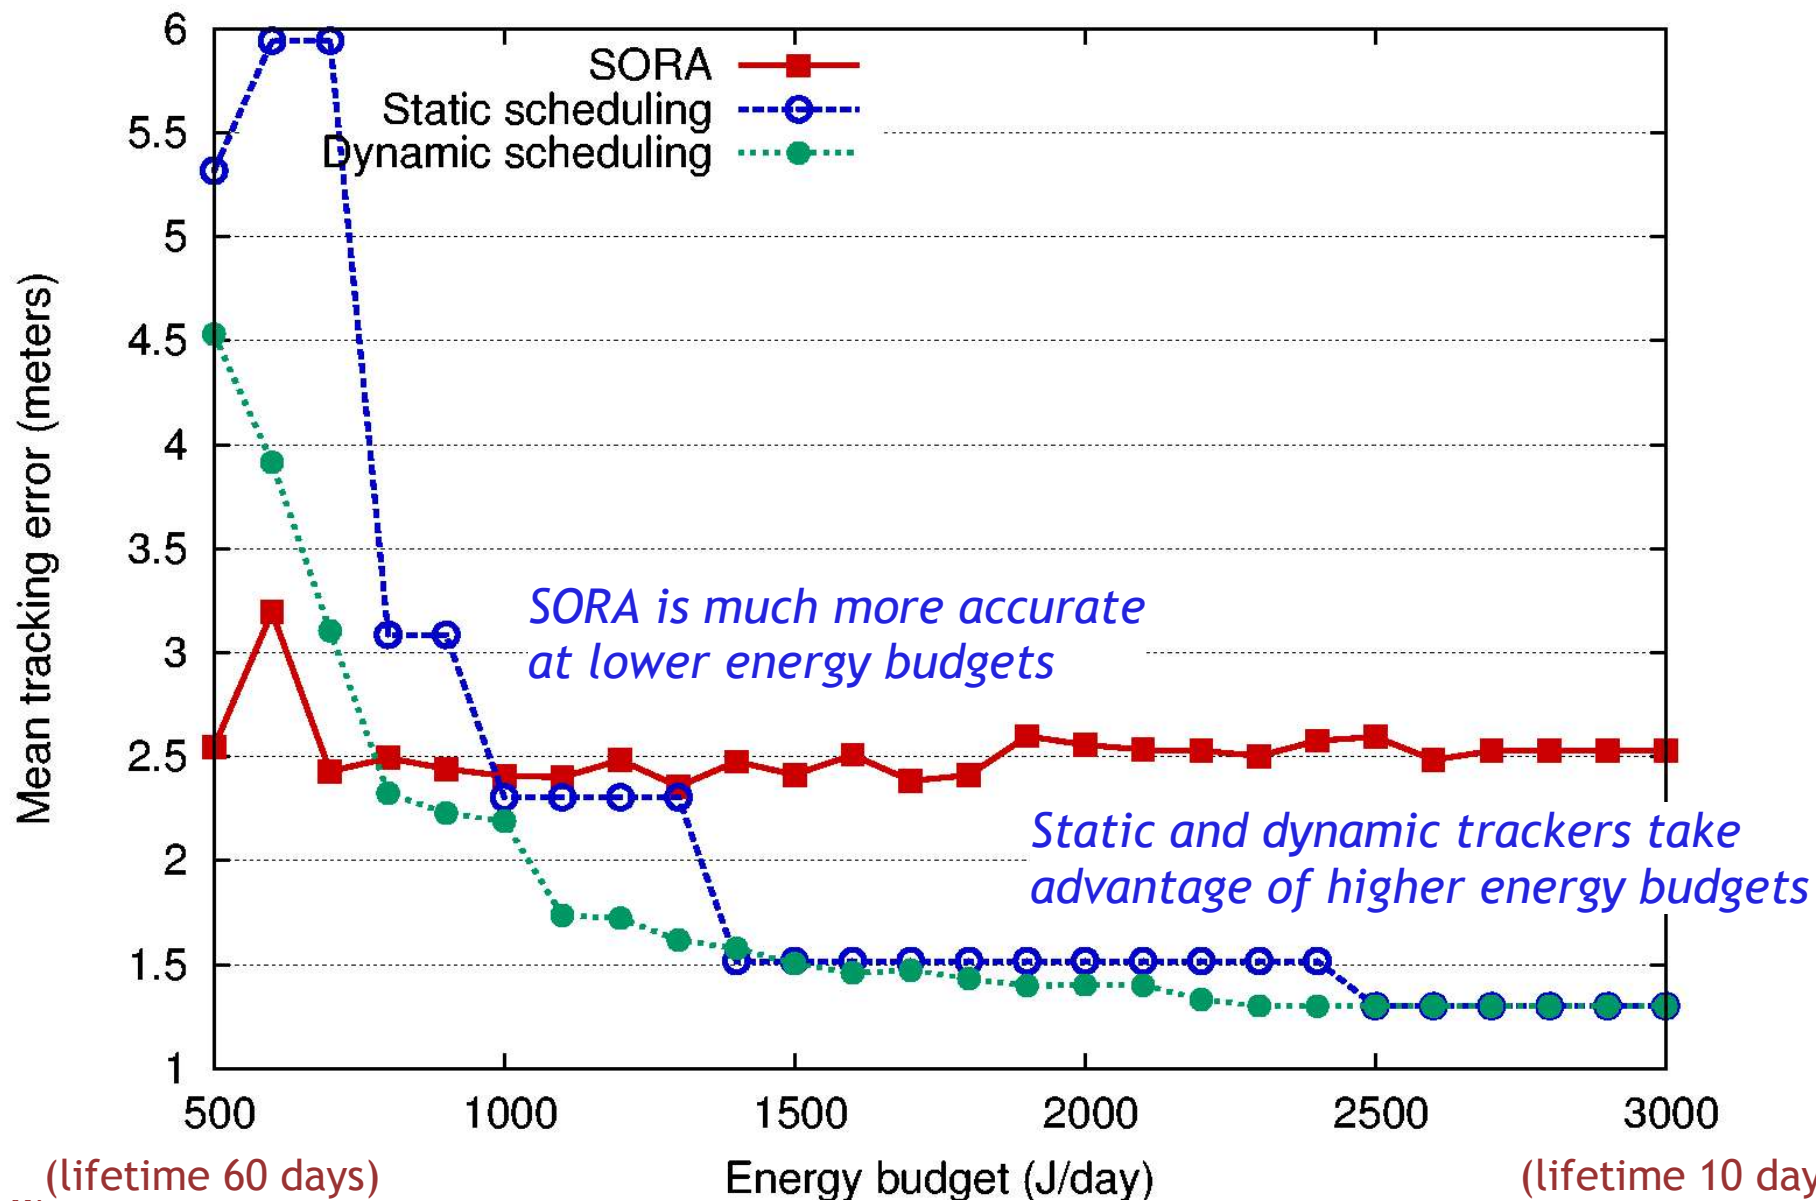
# Overall Tracking Accuracy

(energy budget 1000J/day)






# Effect of varying energy budget on accuracy



# Energy Use



# Energy Efficiency

A key goal of SORA is to maximize the *efficiency* of the network:


$$\text{Efficiency} = \frac{\text{Total energy used to collect and route sensor data}}{\text{Total energy used by all nodes}}$$

- We calculate this by tracking all energy use that resulted in a position estimate arriving at the base station.

Higher efficiency implies less wasted energy

- No realistic system can be 100% efficient
- Wasted energy due to taking bad sensor readings, listening at wrong times, etc.

# Energy Efficiency



# What about varying prices?

We did extensive measurements with different rewards for each action.

Surprisingly, had little effect on tracking accuracy or energy use!

Observation:

- Prices only “matter” when a node has multiple actions with non-zero utility
- But a node can usually take only one action at a time!
  - *At least, this is the case in our tracking application.*

Most of the behavior in SORA is dictated by the learning process, not the choice of reward prices.

# Also in the paper...

## Experiments varying the learning parameters $\epsilon$ and $\alpha$

- These impact the learning behavior and energy efficiency

## Experiments using heterogeneous energy budgets

- Give some nodes a large energy budget (e.g., connected to mains power)

## Experiments with non-uniform reward settings

- Configure some nodes as “routers” and others as “sensors”

# Future Directions

## Allow nodes to reason about future opportunities for profit

- Current scheme very myopic: Nodes always pick most profitable action
- Would like to price valuable *sequences* of actions
  - *e.g., Must sample multiple times before aggregating*

## Extend model to allocate resources across multiple users

- Each network user can pay for different sets of actions
- Use *equilibrium pricing* to seek Pareto optimal resource allocations

## Use reward settings to retask sensor nodes on the fly

- e.g., Nodes on the edge of the network can act as “sentries” detecting vehicle arrival
- Interior nodes can stay dormant
- When sentry detects vehicle, floods a new reward vector to retask interior nodes

# Conclusions

Sensor networks need new tools for managing resources

- Energy and bandwidth are very constrained
- Manual scheduling and allocation is difficult to get right

Our approach: Self-Organizing Resource Allocation (SORA)

- Decentralized, adaptive scheduling of individual node operations
- Nodes use reinforcement learning to tune their behavior over time

SORA achieves:

- High tracking accuracy (nearly as good as “static” scheduling techniques)
- Very low energy usage (nodes learn when to activate on short time scales)
- High energy efficiency (little wasted energy taking useless actions)

<http://www.eecs.harvard.edu/~mdw/proj/mp>

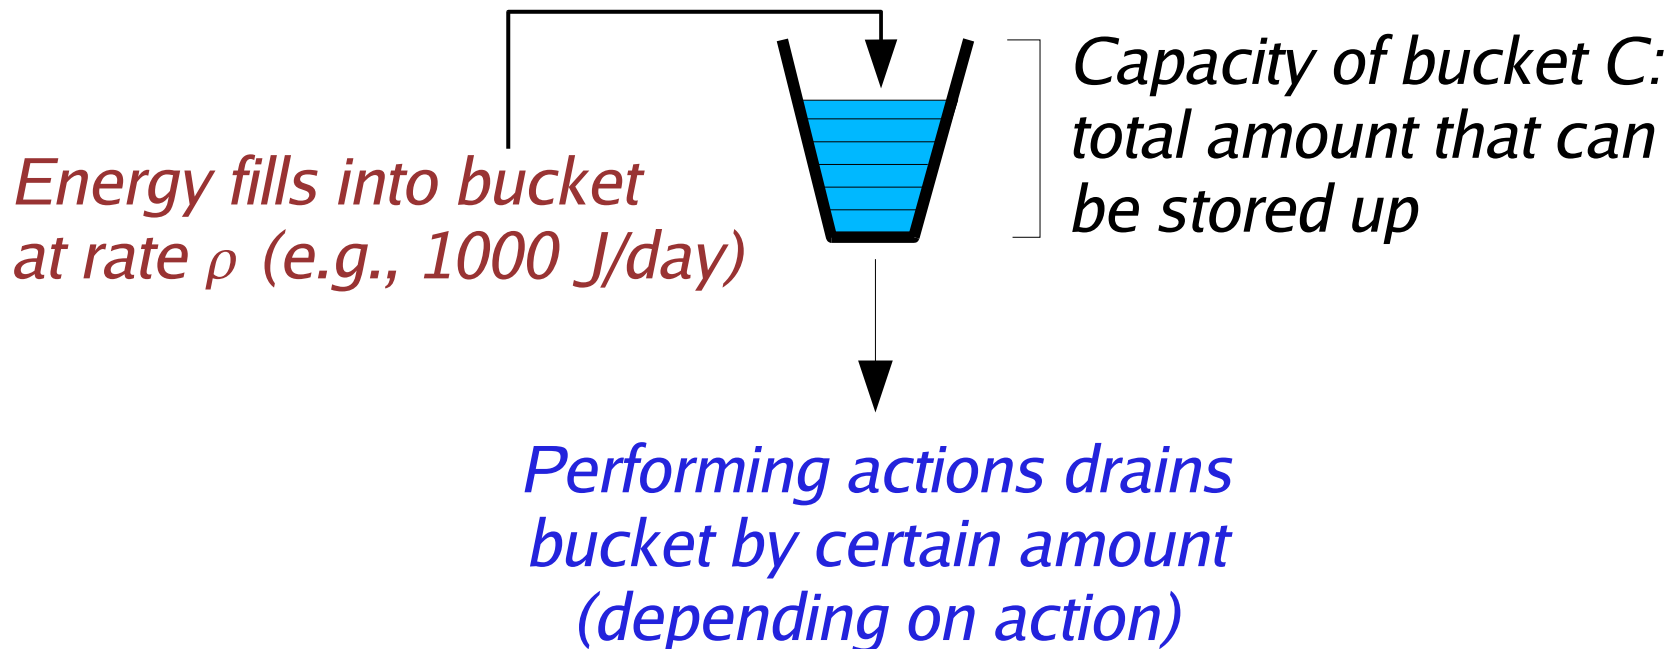


# Energy Budget


Most important constraint on node operation is energy

We model the energy budget for each node as a *token bucket*

- Rate of bucket fill determines average rate of energy use
- Capacity of bucket bounds “burst size”



# Effect of varying exploration probability



$\epsilon$  parameter determines how often node selects a random action

- Low  $\epsilon$ : Node usually chooses highest-utility action
- High  $\epsilon$ : Allows node to find new profit faster

*Low  $\epsilon$ : most energy wasted taking high-utility (but not useful!) actions*

*High  $\epsilon$ : most energy wasted exploring the action space*

*Best setting seems to be somewhere in the middle*

# Actions and energy cost

Nodes can select from four actions:

*Sample* the magnetometer (84  $\mu$ J)

- Results in sample value that scales with distance to vehicle
- Cannot detect vehicle if more than 11 m away

*Listen* for incoming radio messages (5.9 mJ)

*Aggregate* multiple accumulated readings (84  $\mu$ J)

- Computes partial centroid of accumulated values

*Transmit* a message towards the base station (2.4 mJ)

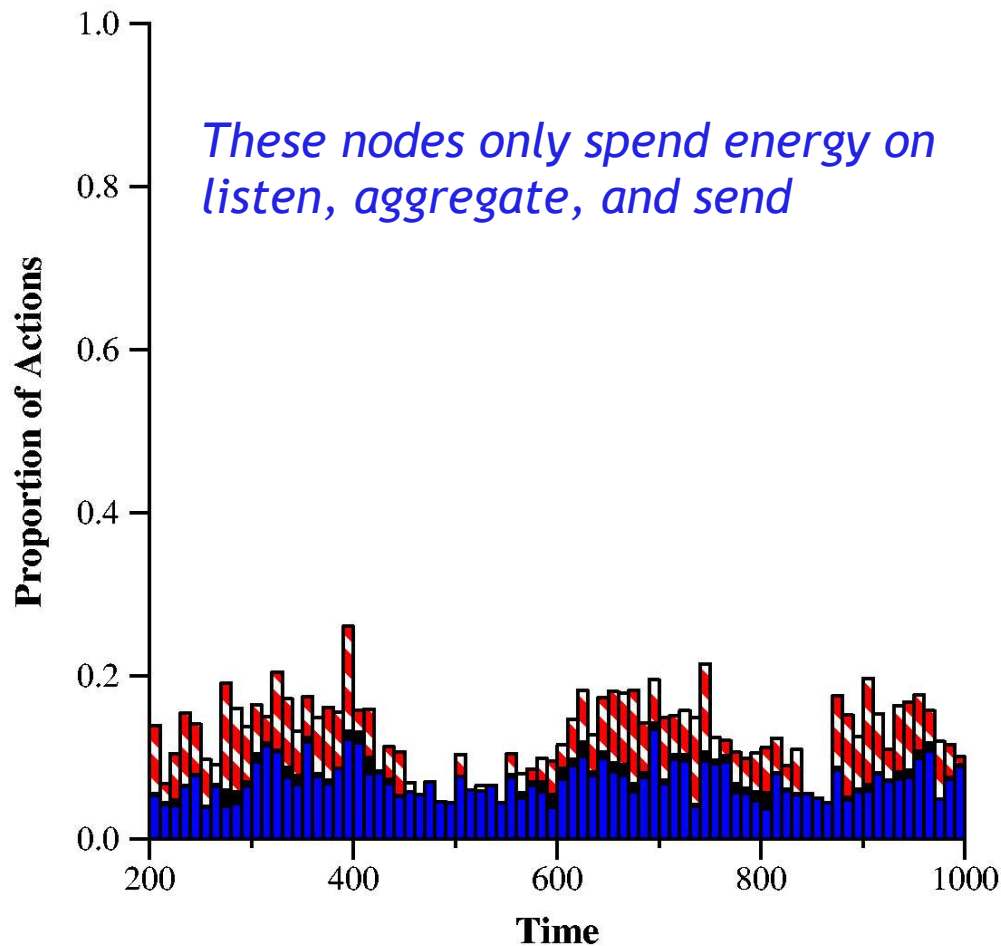
- Uses GPSR routing
- Any node closer to the base that is currently *listening* will receive the message

# Effect of Prices on Action Choice

## “Routers”

(20% of nodes)

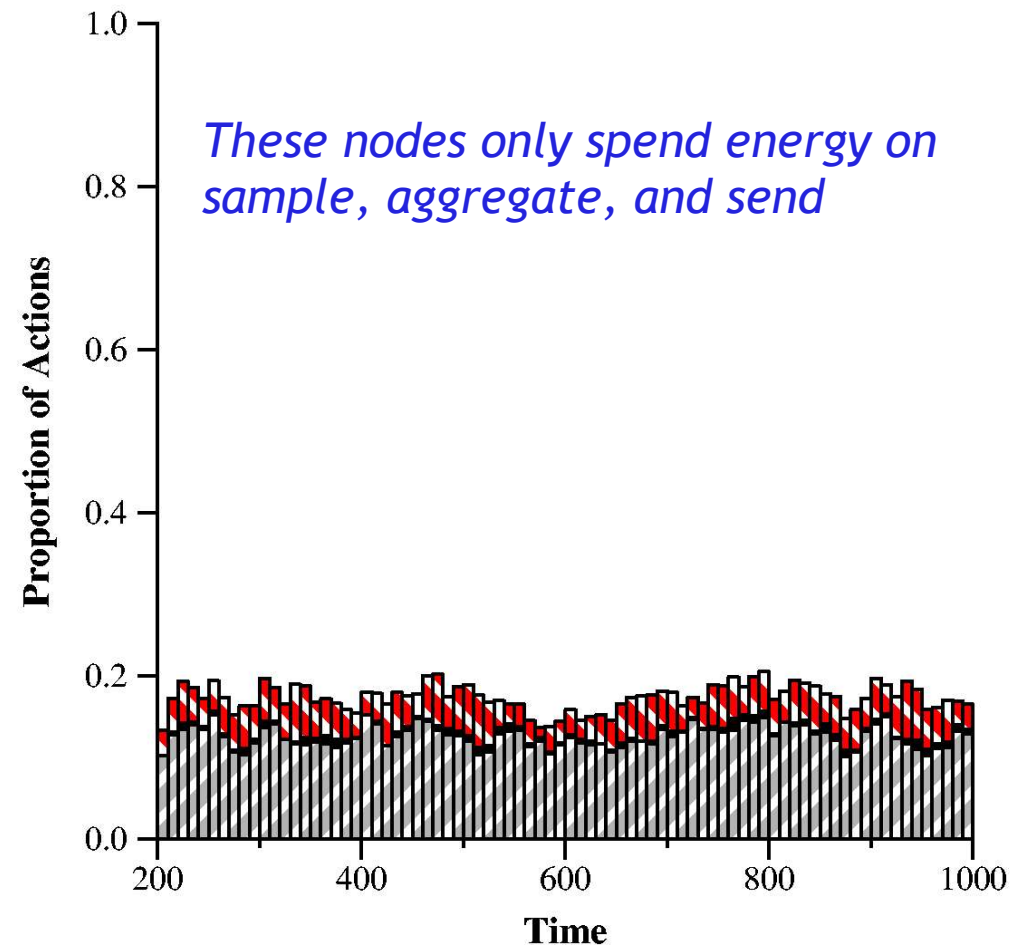
Listen Sample Agg Send



## “Sensors”

(remaining 80%)


Listen Sample Agg Send



# Price Propagation

First step: Flood prices for each good to the network

- Several efficient protocols for this (e.g., Trickle)
- Can readily update prices on the fly



# Comparison to Alternatives


## Hoods [Whitehouse et al., MobiSys'04]

- Programming abstraction for neighborhood-based communication

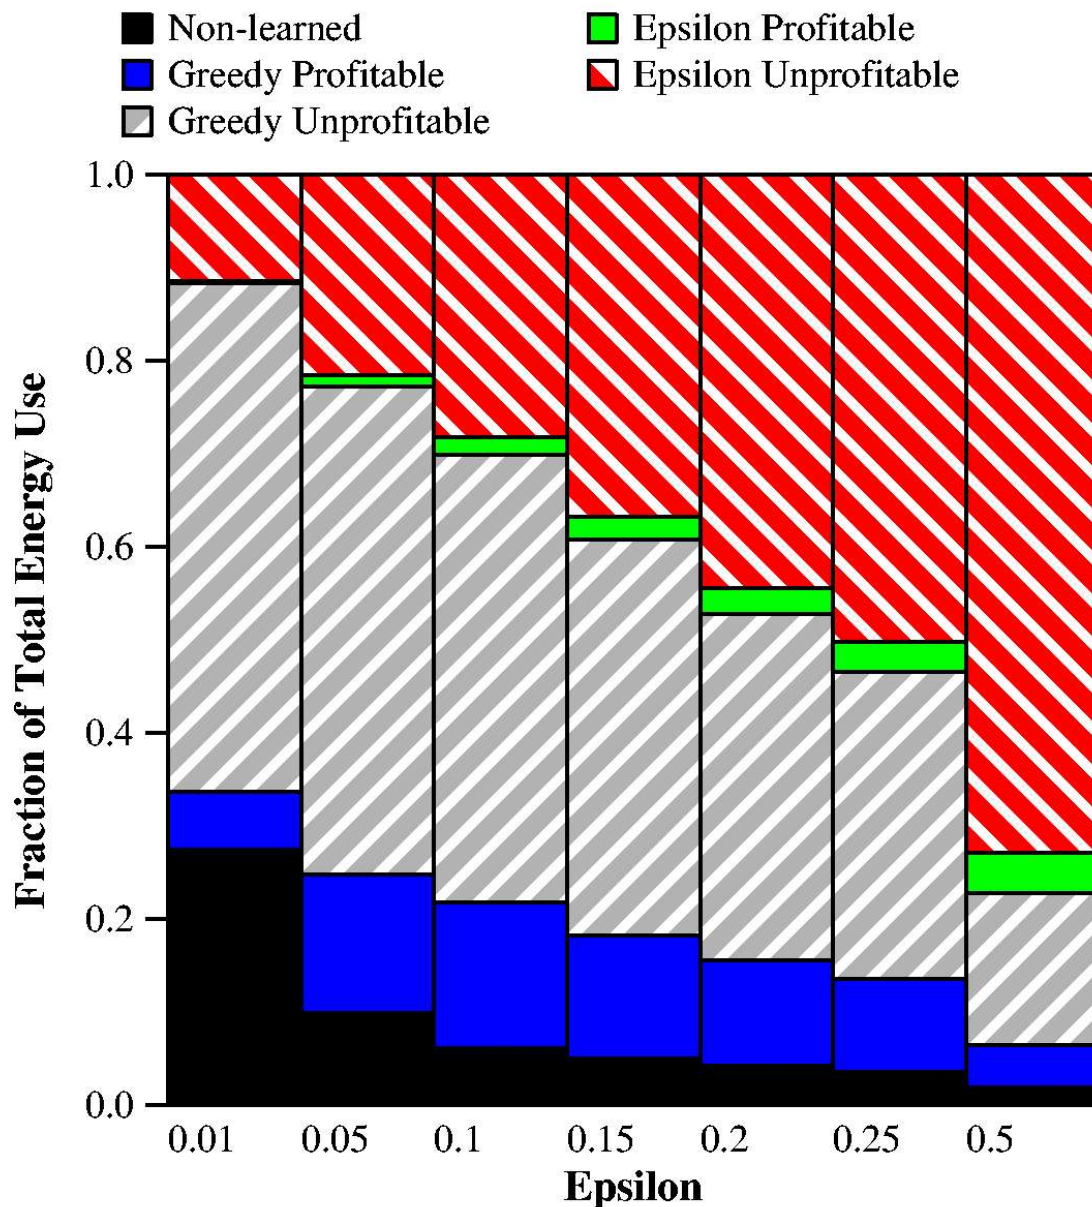
## Hoods implements a different approach to tracking:

- Nodes broadcast sensor values to neighborhood
- “Leader” node aggregates data and sends position estimate to base station
- We found that this is less accurate than the SORA, static, and dynamic trackers

# Actions taken over time



# Effect of varying exploration probability



$\epsilon$  parameter determines how often node selects a random action

- Low  $\epsilon$ : Node usually chooses highest-utility action
- High  $\epsilon$ : Allows node to find new profit faster

*Low  $\epsilon$ : most energy wasted taking high-utility (but not useful!) actions*

*High  $\epsilon$ : most energy wasted exploring the action space*

*Best setting seems to be somewhere in the middle*